

# IMAGE STYLIZATION BY AI

PRESENTED BY GOURAV CHOUDHARY - 13, ISHANK VARSHNEY - 15,  
ISHITA JAIN - 16 AND VISHVENDRA SINGH - 56

# WHAT IS IMAGE STYLIZATION BY AI?

Image stylization is an optimization technique used to take two images—a content image and a style reference image (such as an artwork by a famous painter)—and blend them together so the output image looks like the content image, but “painted” in the style of the style reference image.



# IMAGE STYLIZATION



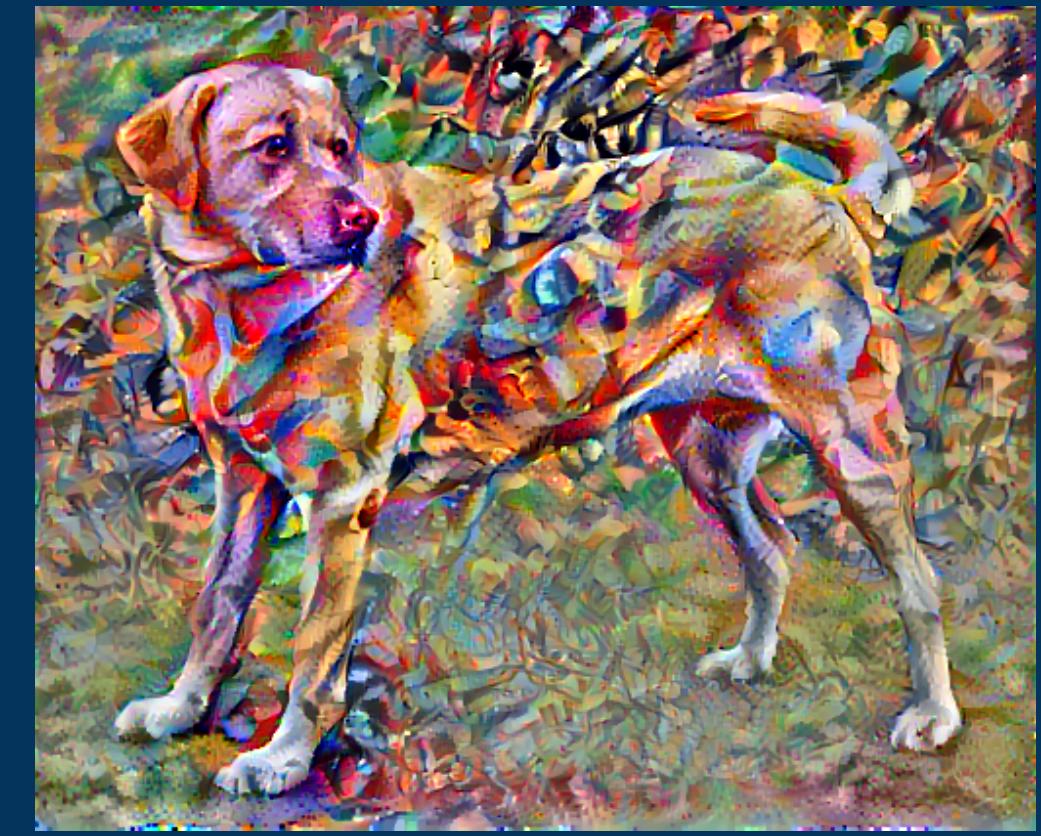
Content Image

+



Style Image

=



Result

# EXAMPLES



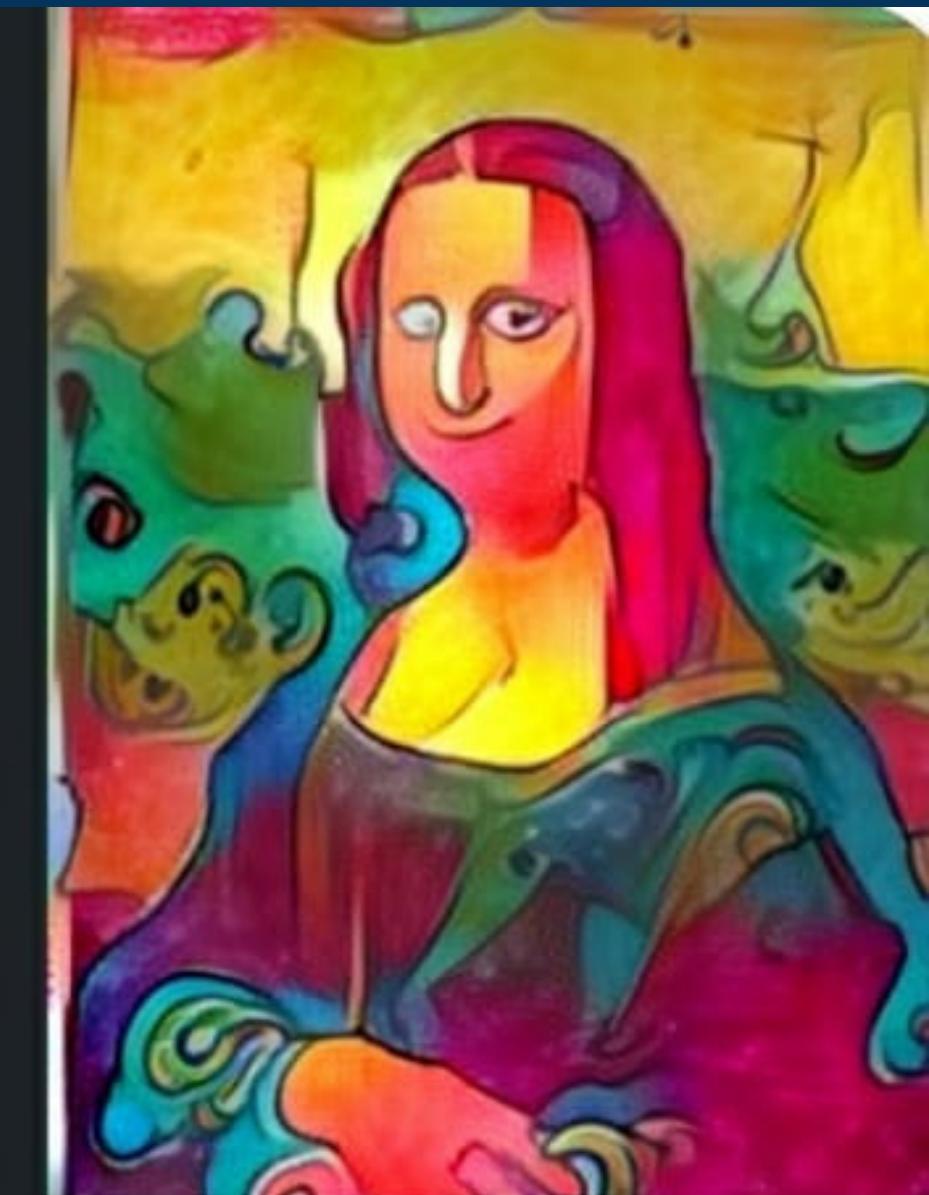
*Content Image*

+



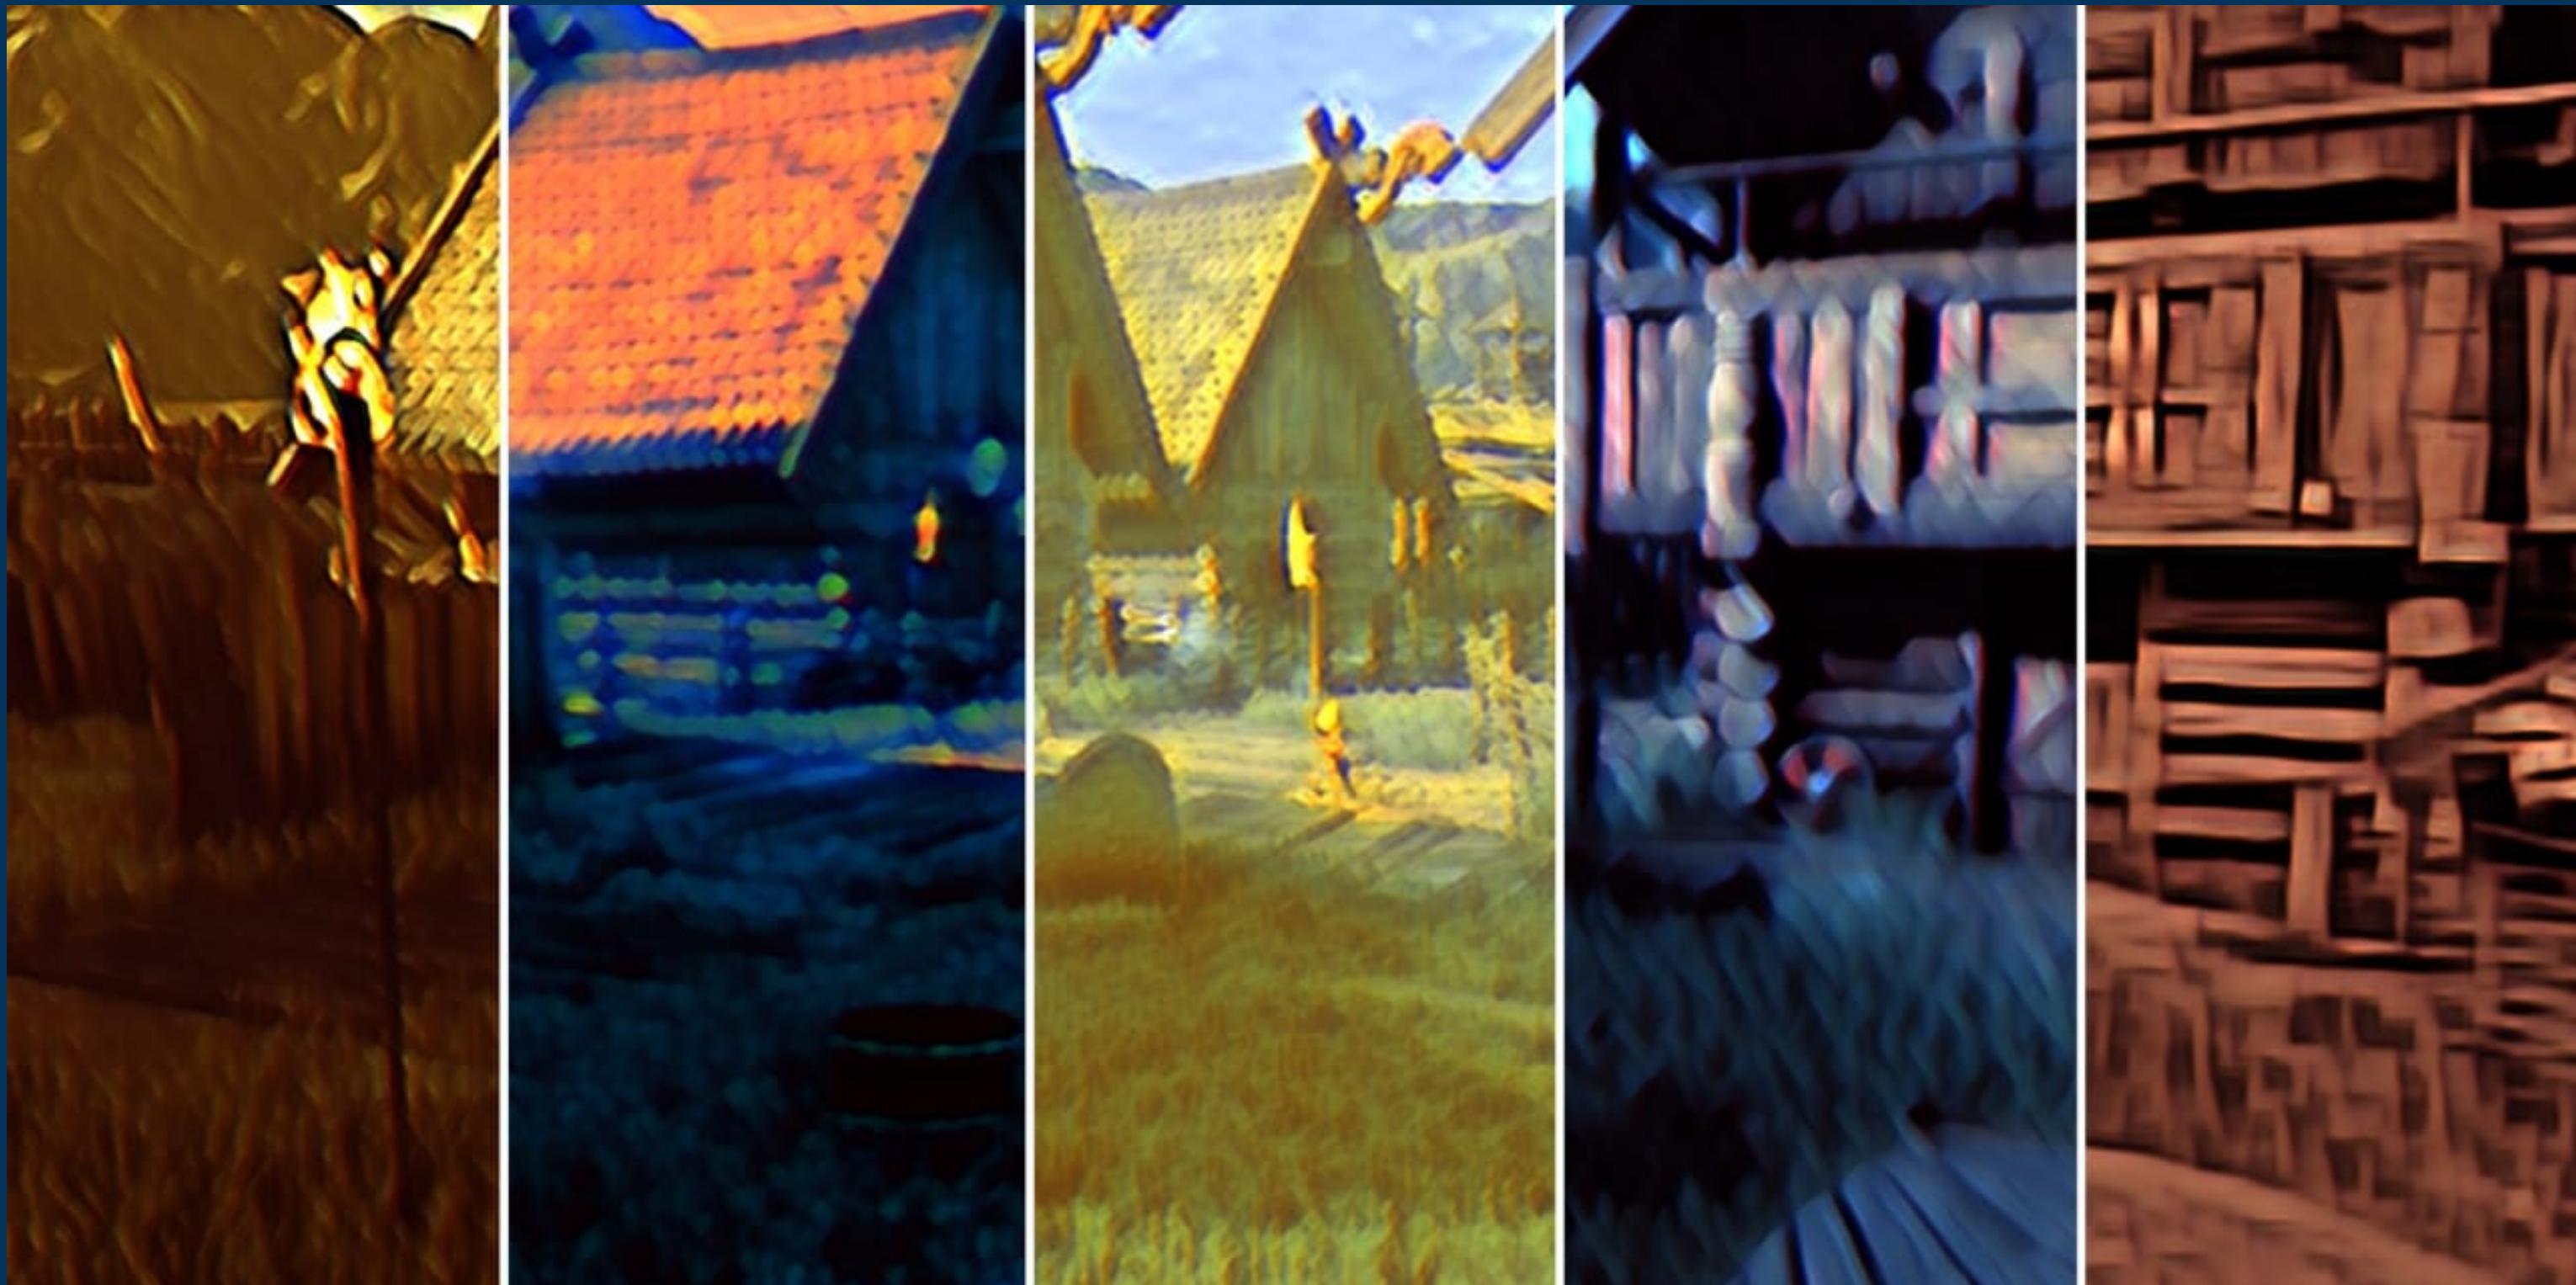
*Style Image*

=



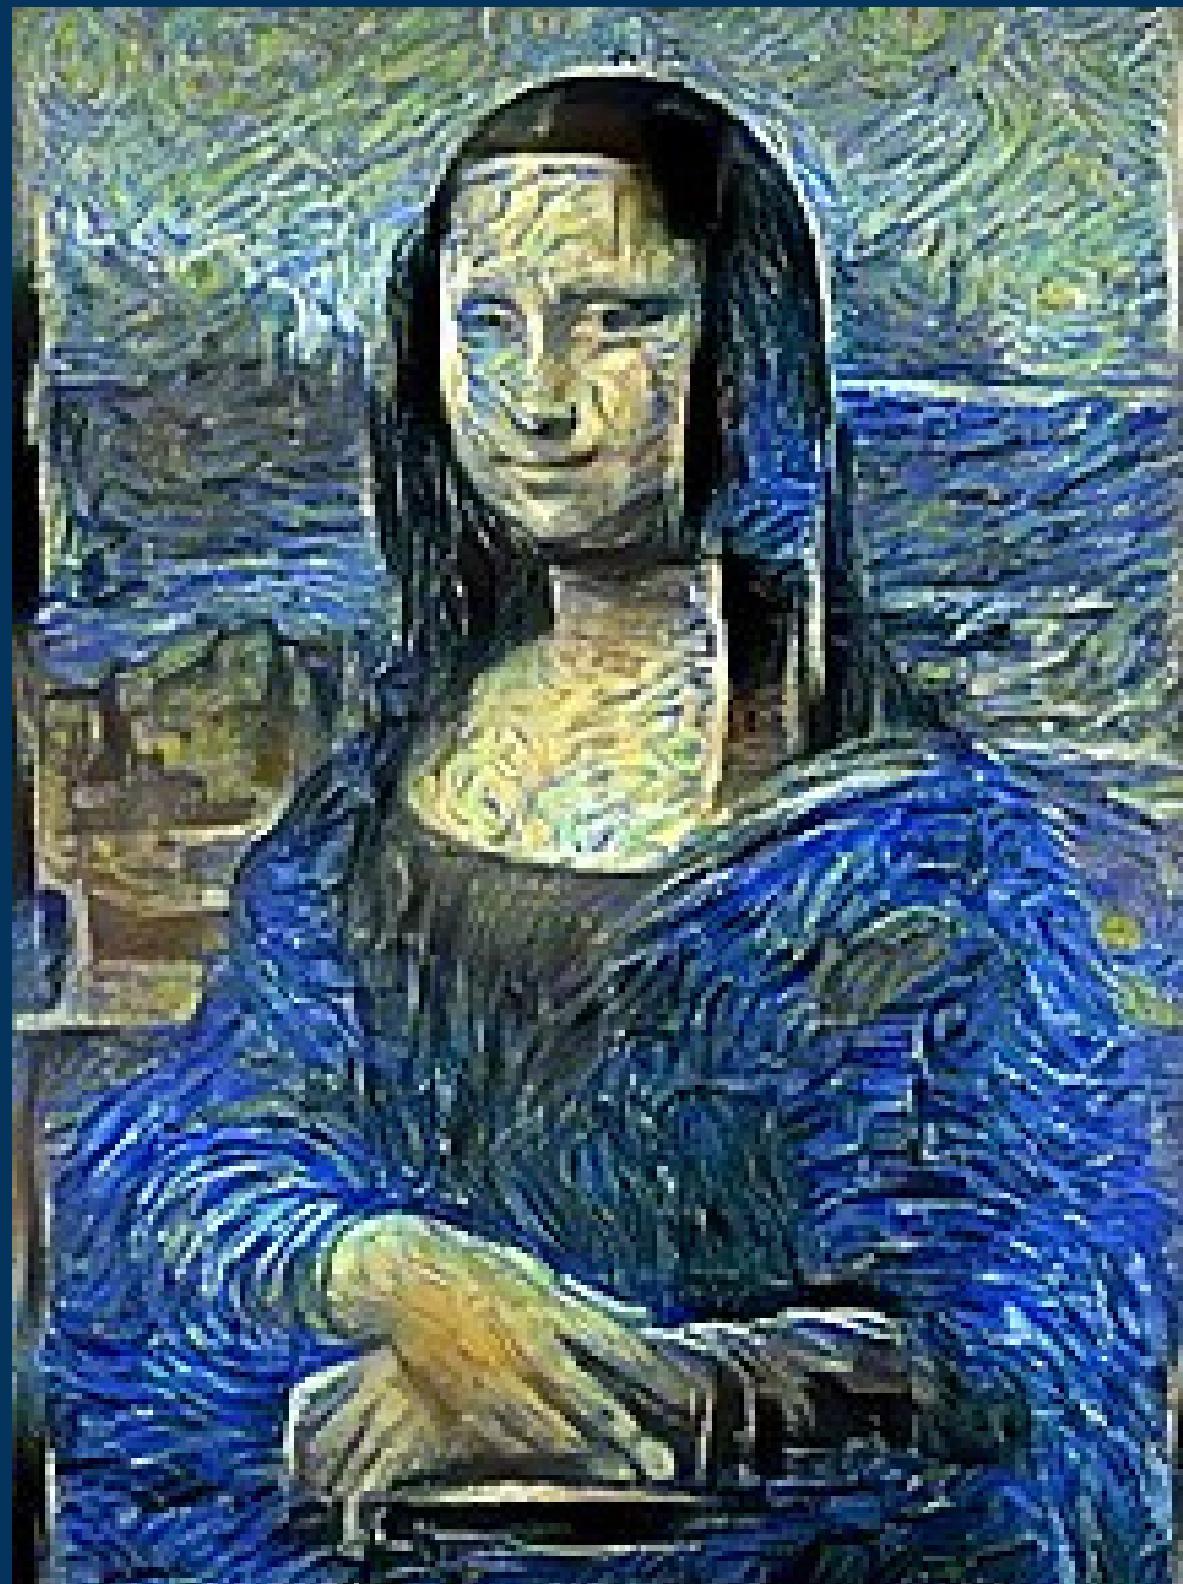
*Neural Style Transfer*

# EXAMPLES



# REAL LIFE APPLICATIONS

WHY IS IT IMPORTANT TO US?



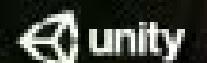
## ARTISTIC SATISFACTION

Not everyone is born an artist. Some are more adept at language or other tactile tasks. But with recent advances in technologies like style transfer, almost anyone can enjoy the pleasure that comes along with creating and sharing an artistic masterpiece..

## LIVELY ENVIRONMENT

Image stylization helps us make world a beautiful place. It inspires creativity in the art world.

The lens created using this can help people in viewing their surroundings in a new and creative light..



unity [unity3d.com/book-of-the-dead](http://unity3d.com/book-of-the-dead)

# ART EXHIBITIONS FOR AI

Image stylization can be used by artist to exhibit the art pieces in art galleries that display and exhibition the art pieces that are designed and created by AI..



## APPLICATION IN GAMING

Image stylization can be used by game developers to create some attractive and creative effects for various functions in the game. .



# LIBRARIES THAT WILL BE USED

- **TensorFlow:** Used for deep learning and neural networks. TensorFlow is a free and open-source software library for machine learning and artificial intelligence. It can be used across a range of tasks but has a particular focus on training and inference of deep neural networks.
- **PIL.Image(Python Image Library):** Used. for image processing and loading. Python Imaging Library is a free and open-source additional library for the Python programming language that adds support for opening, manipulating, and saving many different image file formats. It is available for Windows, Mac OS X and Linux



```
import tensorflow as tf  
import PIL.Image
```

## Importing Python Libraries

```
[1] import os
    import tensorflow as tf
    os.environ['TFHUB_MODEL_LOAD_FORMAT'] = 'COMPRESSED'

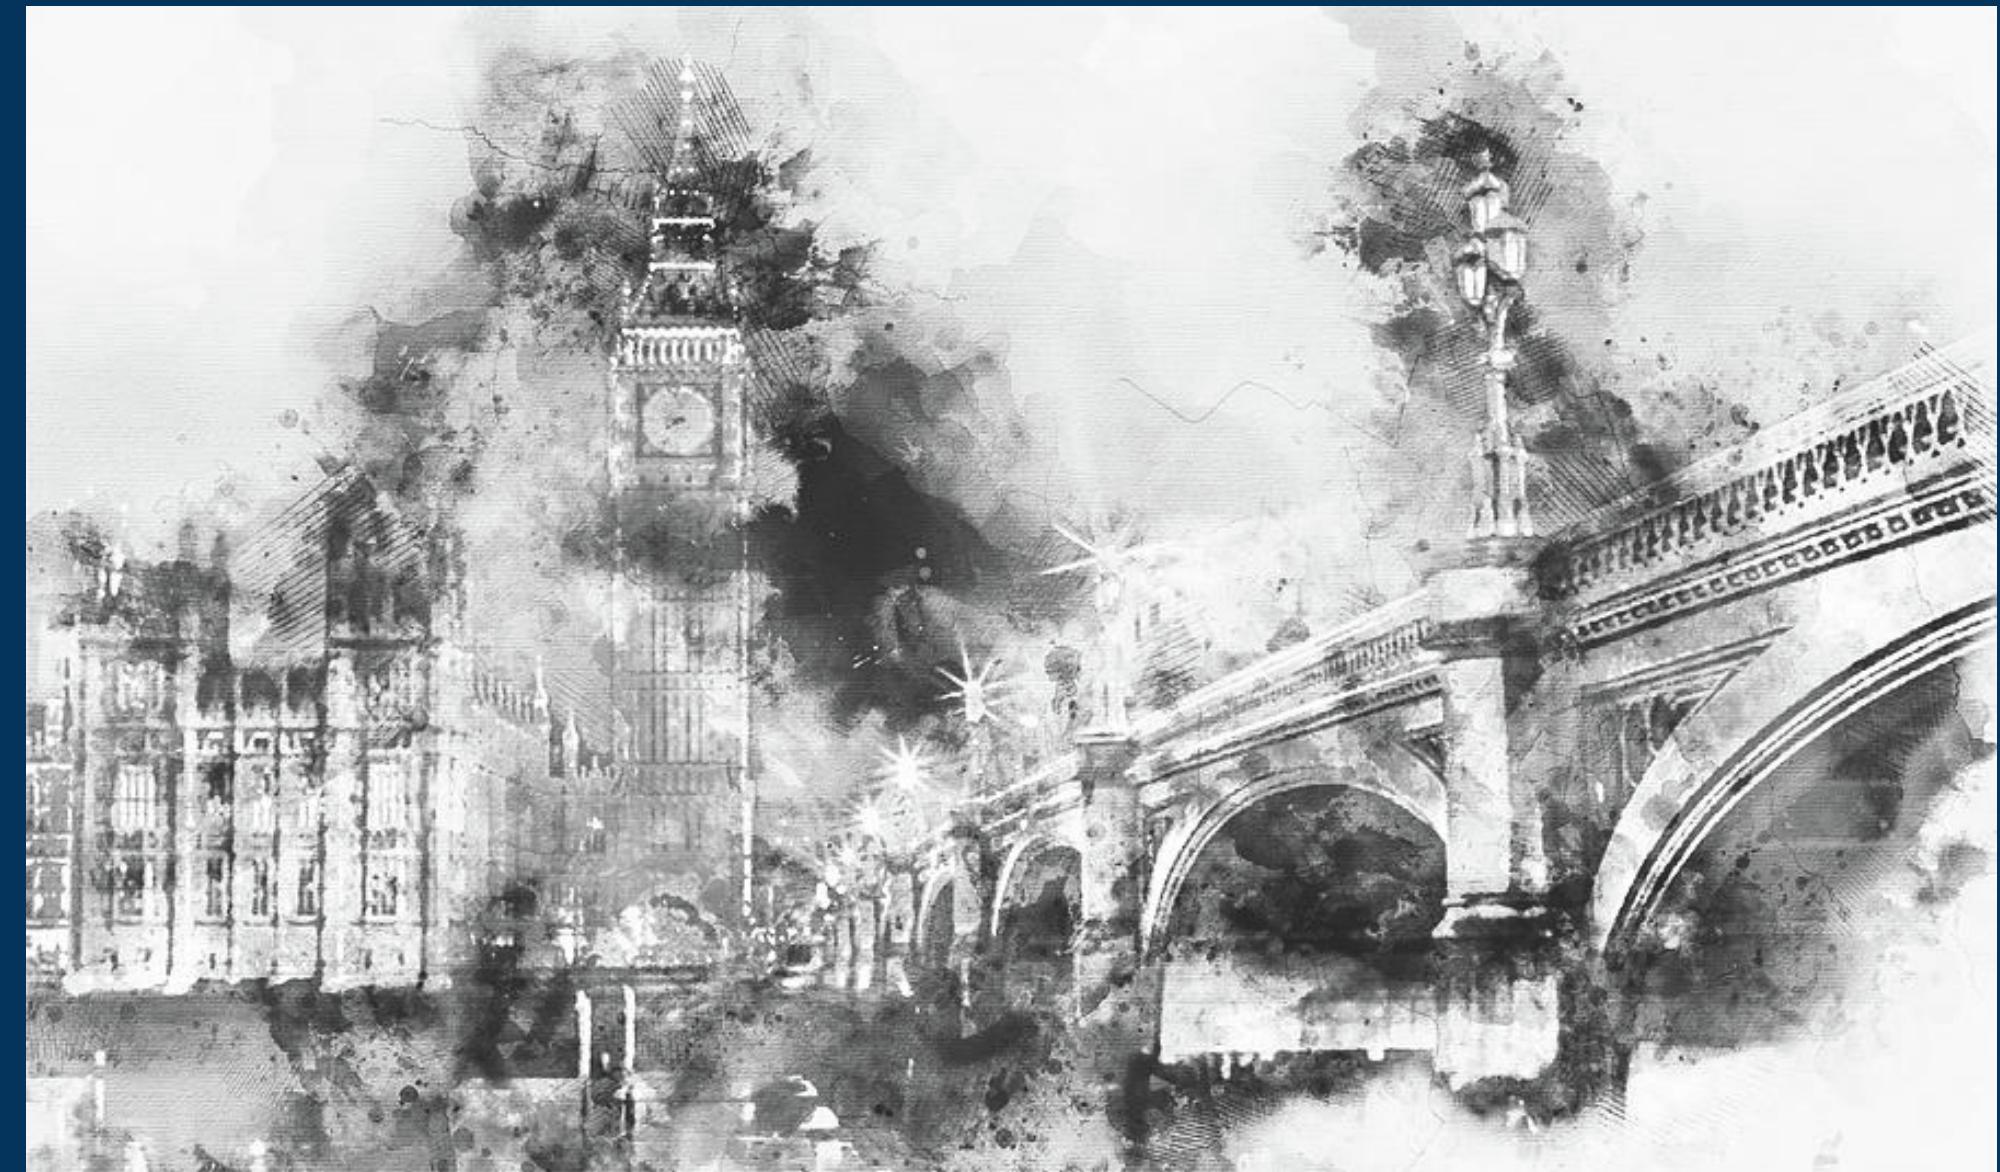
[2] import matplotlib.pyplot as plt
    import matplotlib as mpl
    mpl.rcParams['figure.figsize'] = (12, 12)
    mpl.rcParams['axes.grid'] = False

[3] import numpy as np
    import PIL.Image
```

# IMAGES USED BY OUR MODEL



Content Image



Style Image

```
[7] def tensor_image(tensor):
    tensor = tensor*255
    tensor = np.array(tensor, dtype=np.uint8)
    if np.ndim(tensor)>3:
        if tensor.shape[0] == 1:
            tensor = tensor[0]
    return PIL.Image.fromarray(tensor)
```

```
[49] content_path = "/content/leonardo.jpg"
     style_path = "/content/london.jpg"
```

```
[50] def load_img(path_to_img):
        max_dim = 512
        img = tf.io.read_file(path_to_img)
        img = tf.image.decode_image(img, channels=3)
        img = tf.image.convert_image_dtype(img, tf.float32)

        shape = tf.cast(tf.shape(img)[ :-1 ], tf.float32)
        long_dim = max(shape)
        scale = max_dim / long_dim

        new_shape = tf.cast(shape * scale, tf.int32)

        img = tf.image.resize(img, new_shape)
        img = img[tf.newaxis, :]
        return img
```

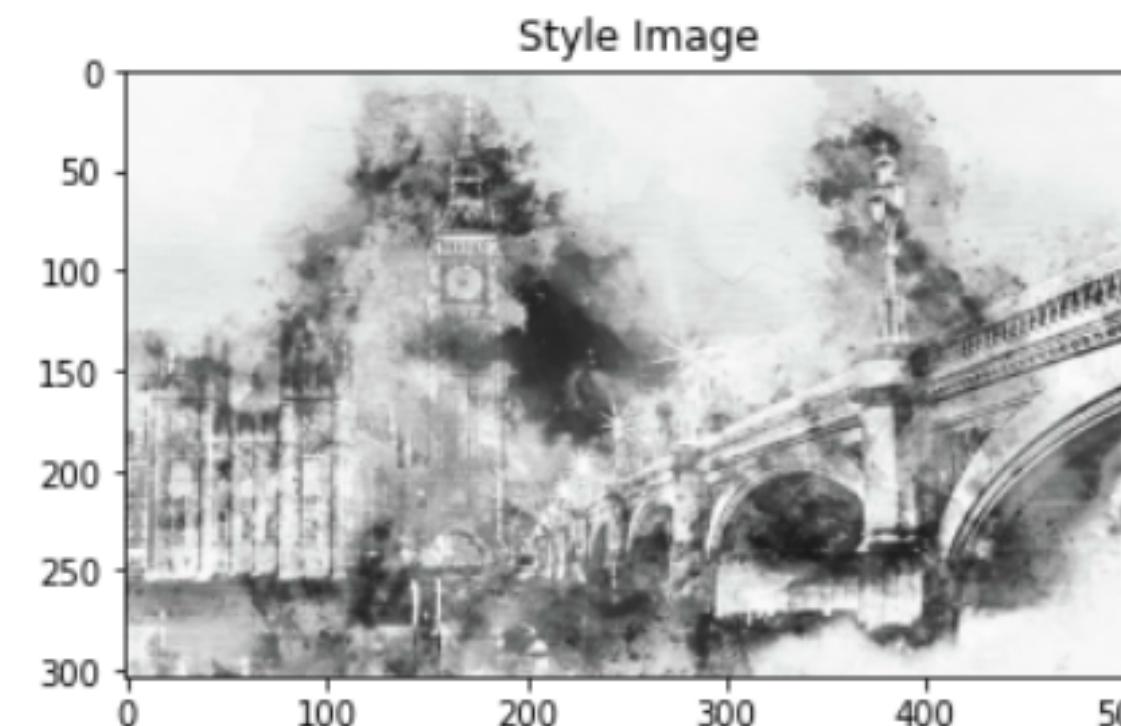
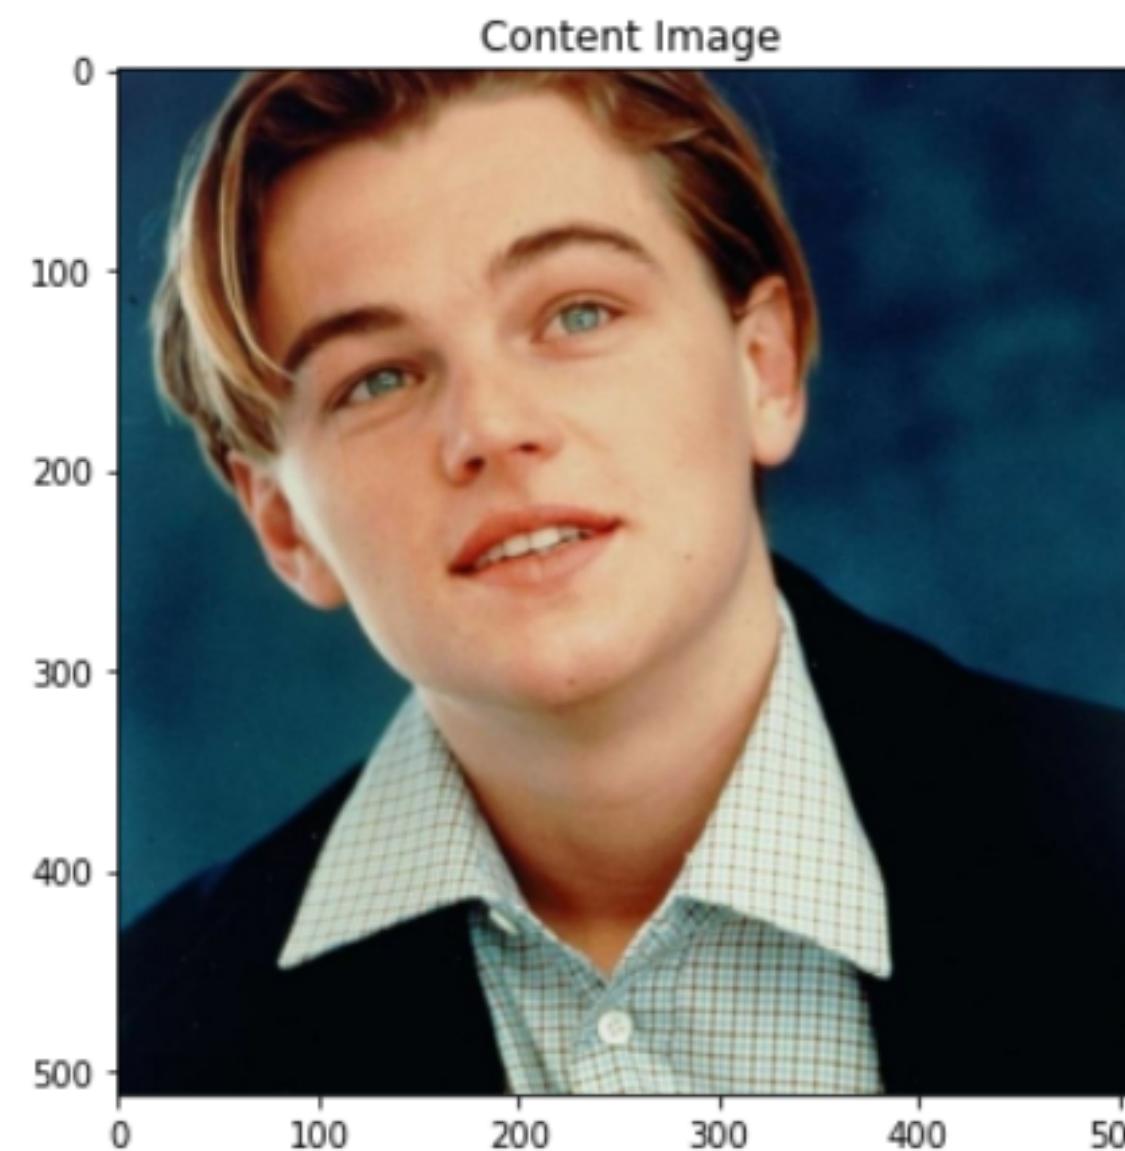
```
[51] def imshow(image, title=None):
        if len(image.shape) > 3:
            image = tf.squeeze(image, axis=0)

        plt.imshow(image)
        if title:
            plt.title(title)
```

```
[52] content_image = load_img(content_path)
      style_image = load_img(style_path)

      plt.subplot(1, 2, 1)
      imshow(content_image, 'Content Image')

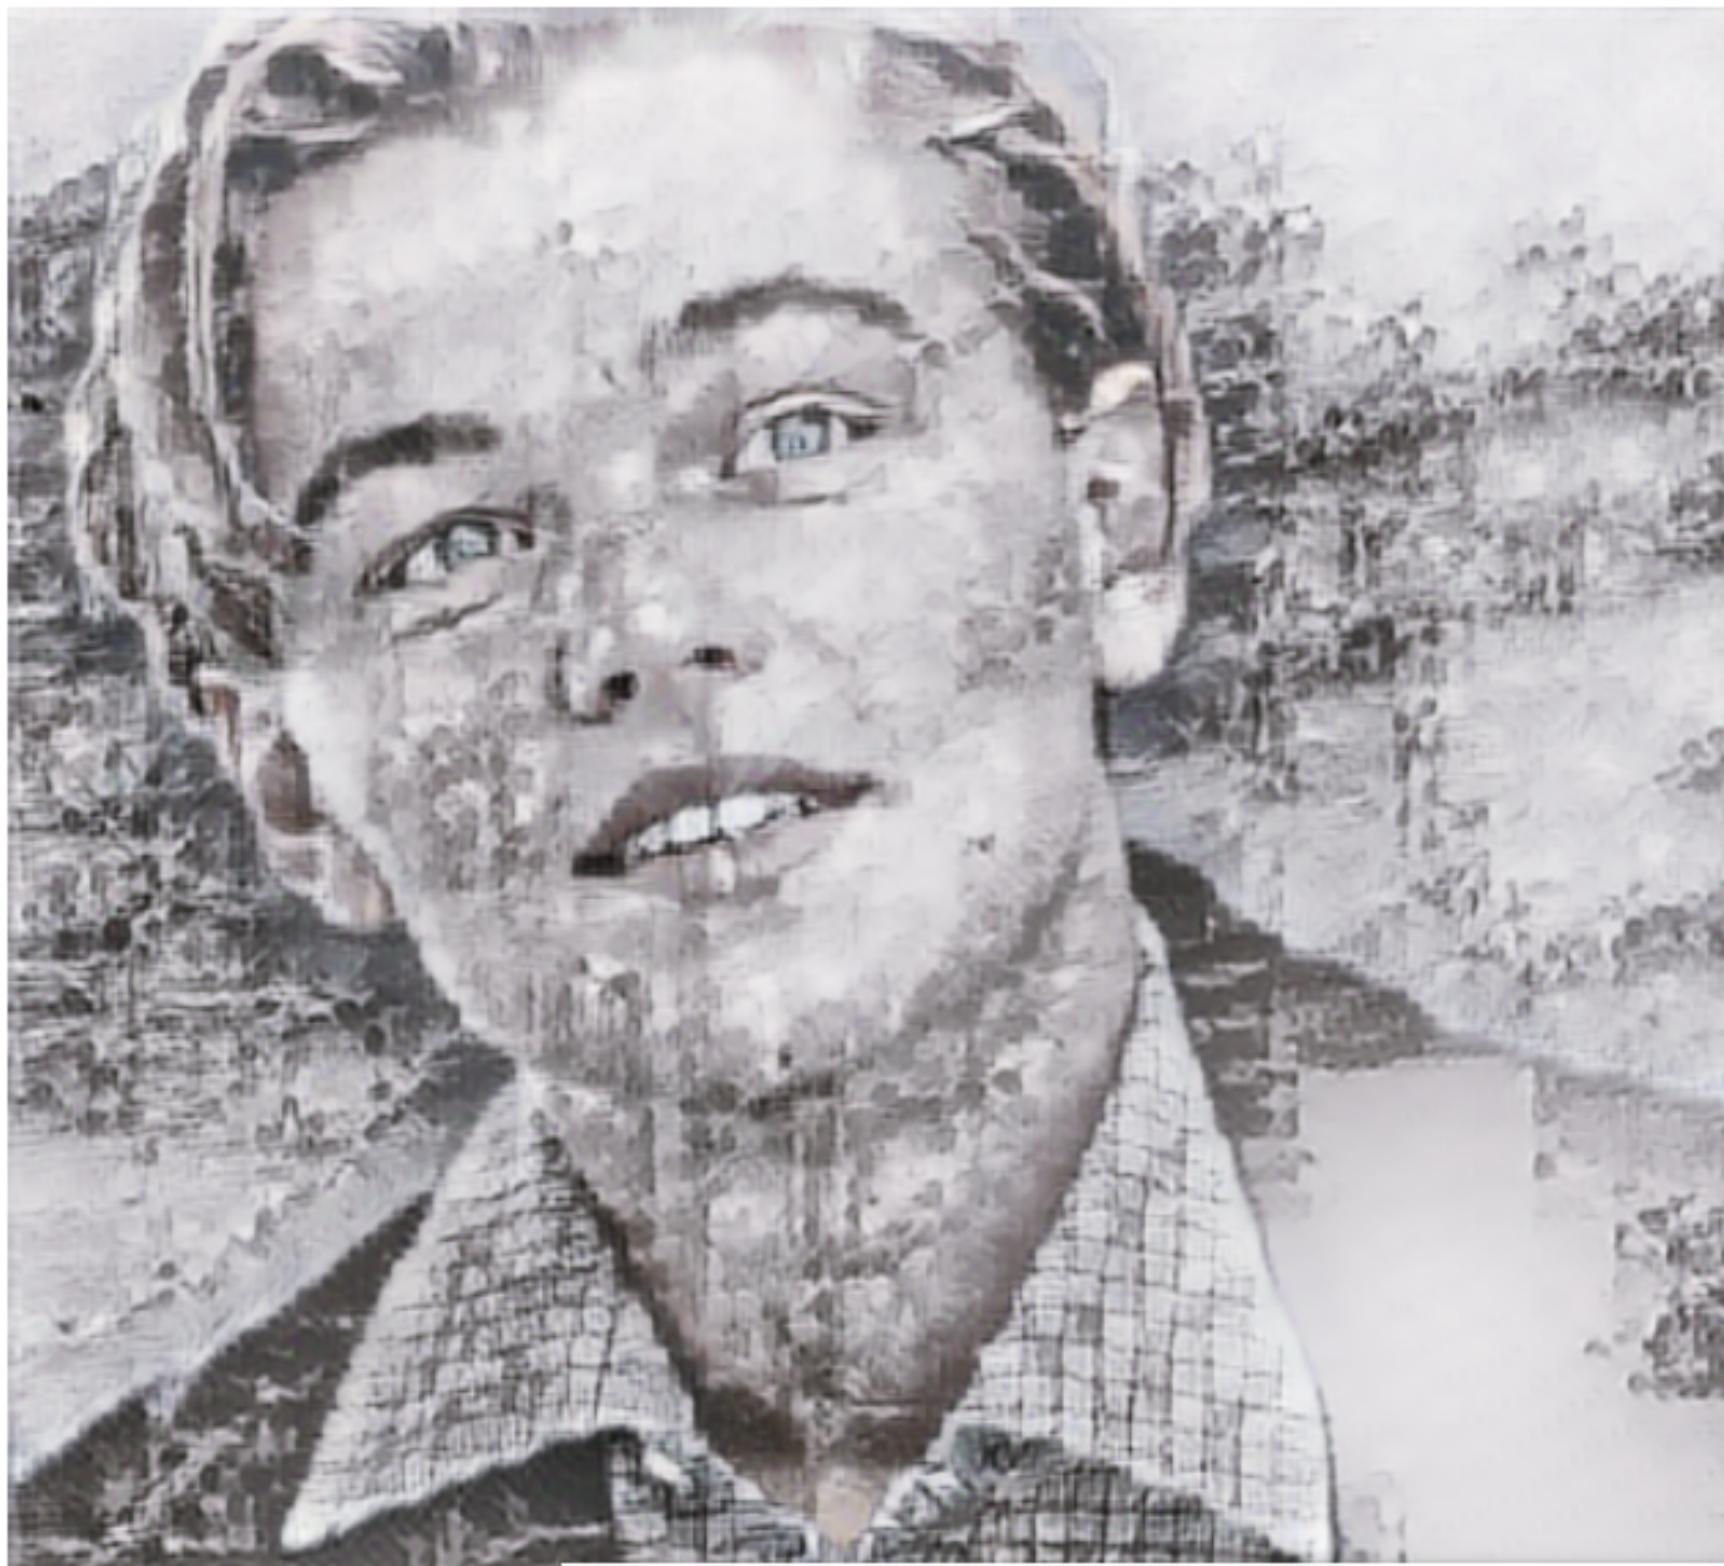
      plt.subplot(1, 2, 2)
      imshow(style_image, 'Style Image')
```



# TENSORFLOW HUB

TensorFlow Hub is a repository of trained machine learning models ready for fine-tuning and deployable anywhere. Reuse trained models like BERT and Faster R-CNN with just a few lines of code. dev repository provides many pre-trained models: text embeddings, image classification models, TF.js/TFLite models and much more.

```
[53] import tensorflow_hub as hub  
hub_model = hub.load('https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-256/2')  
stylized_image = hub_model(tf.constant(content_image), tf.constant(style_image))[0]  
tensor_image(stylized_image)
```

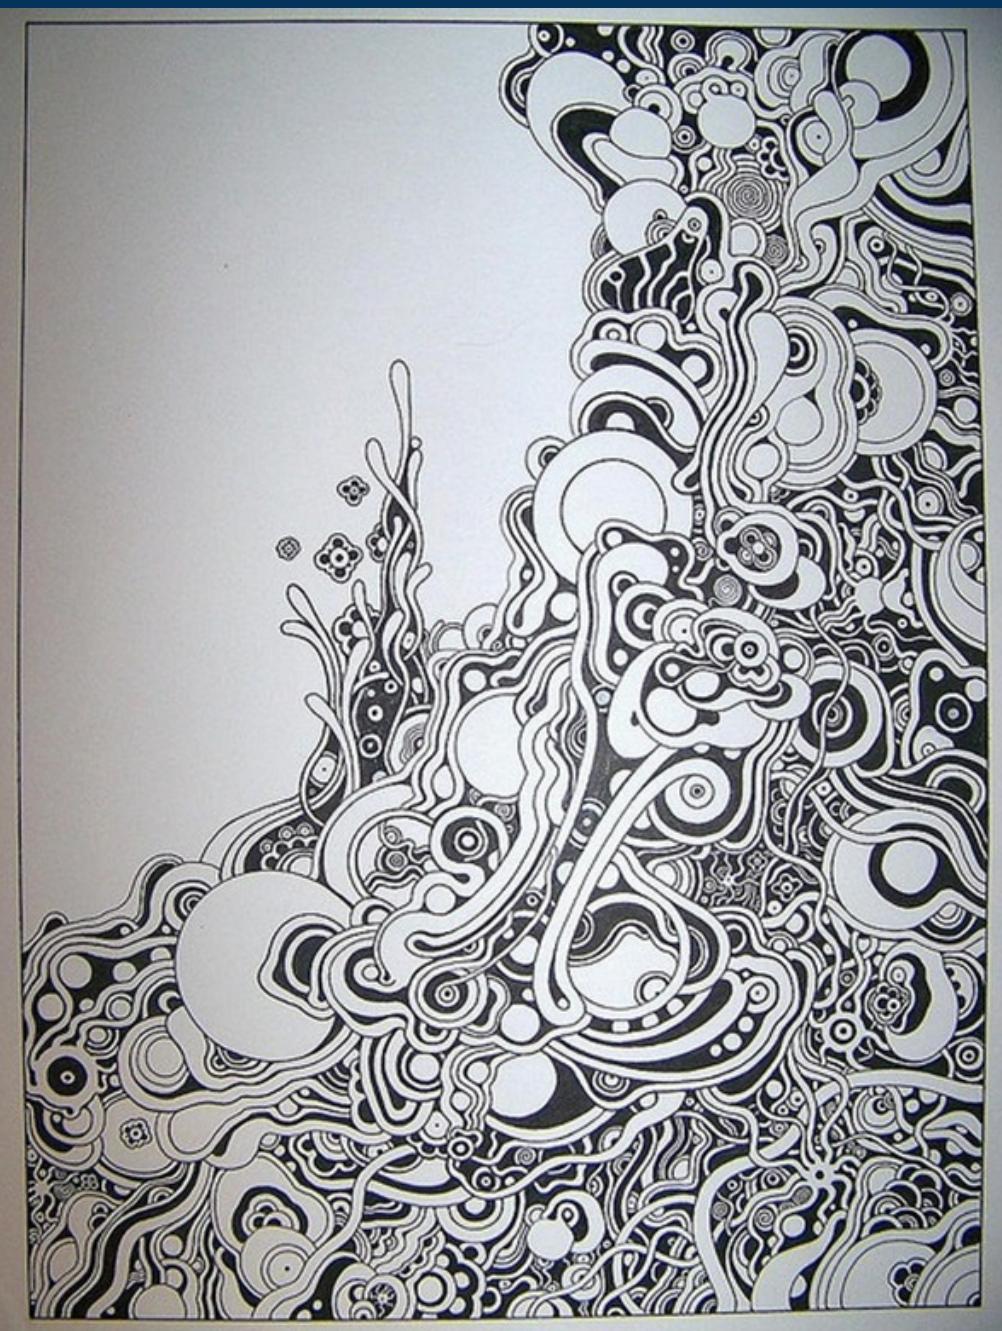


# WHY WE NEED A BETTER MODEL



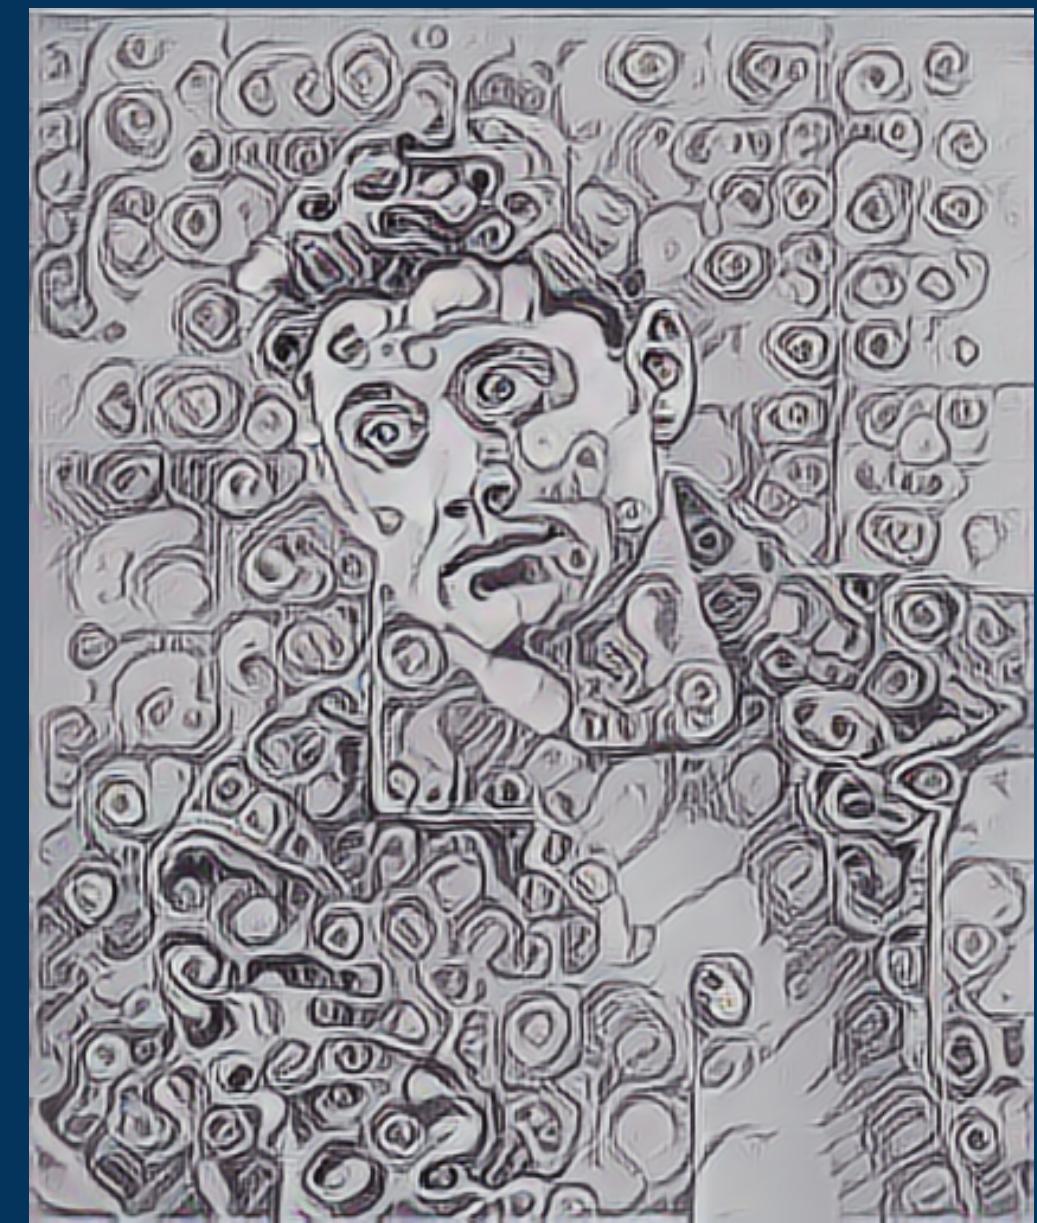
Content Image

+



Style Image

=



Stylised Image

# WHY WE NEED A BETTER MODEL



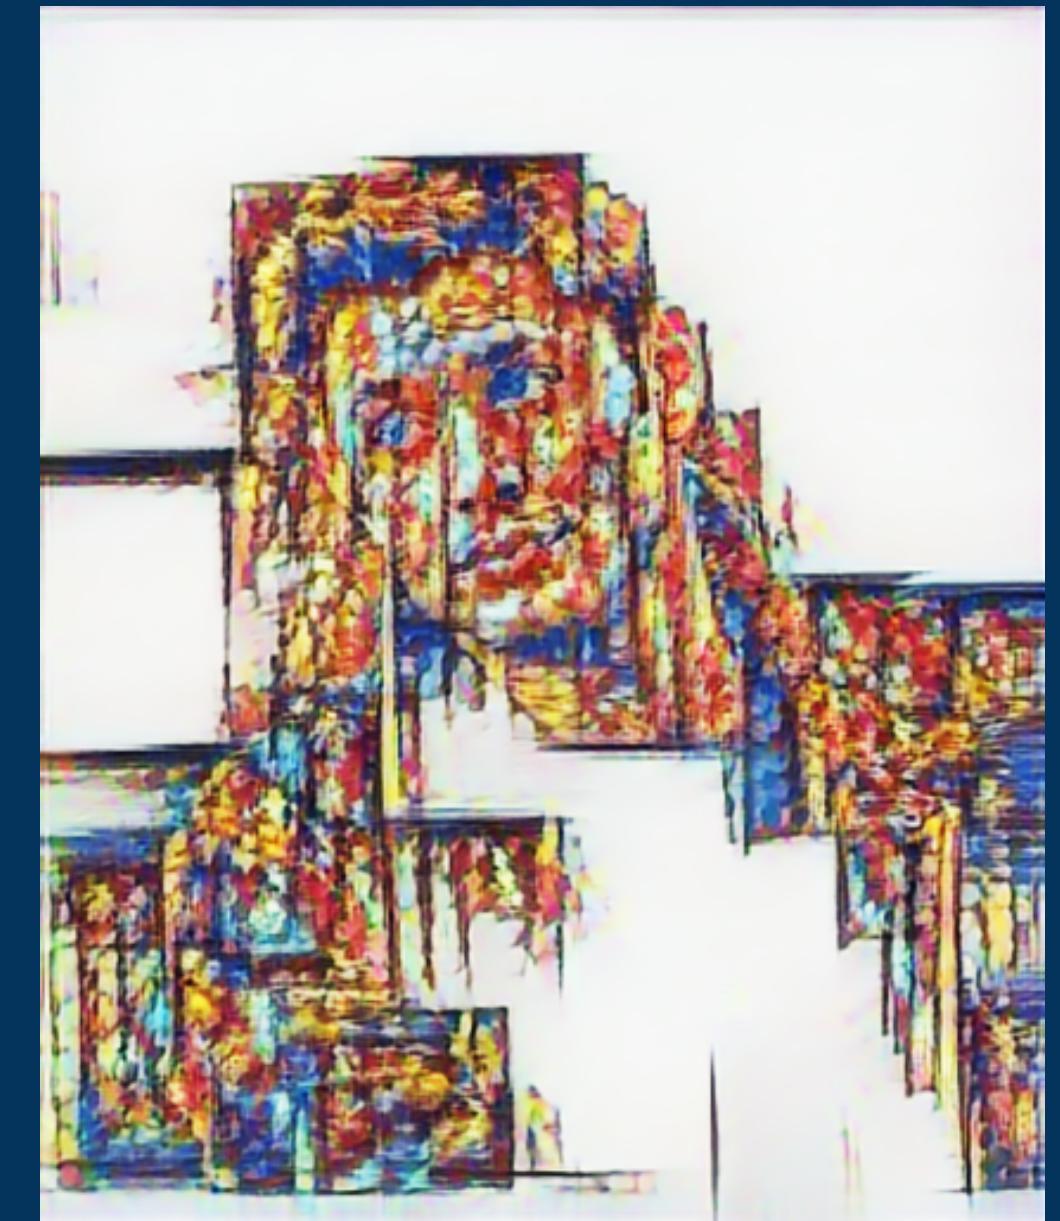
Content Image

+



Style Image

=



Stylised Image

# CONVOLUTIONAL NEURAL NETWORK

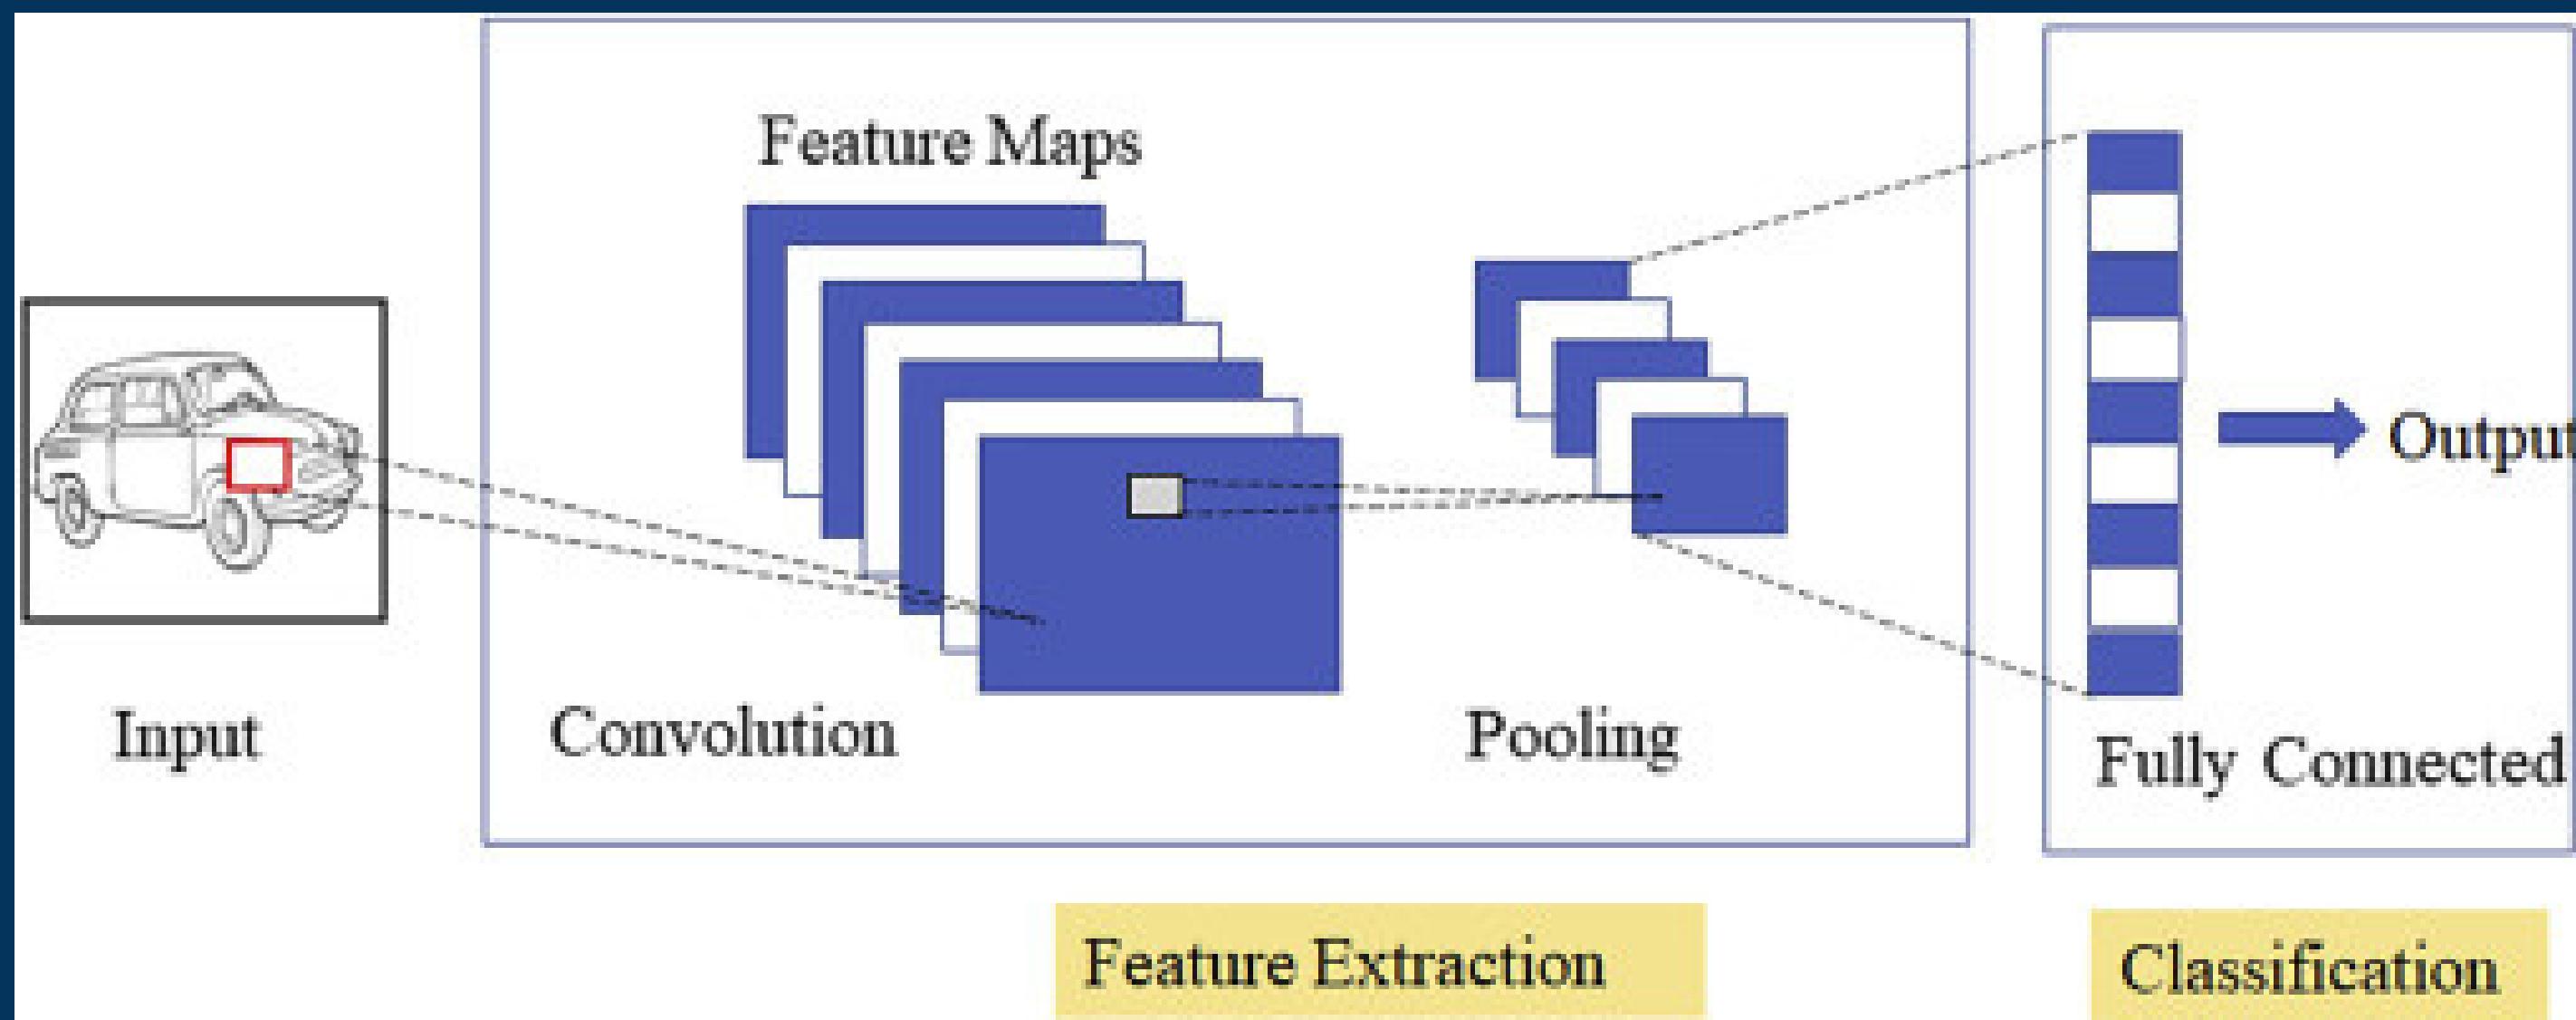
A **Convolutional Neural Network (ConvNet/CNN)** is a Deep Learning algorithm which can take in an input image, assign importance, learnable weights and biases, to various aspects/objects in the image and be able to differentiate one from the other. The pre-processing required in a ConvNet is much lower as compared to other classification algorithms. CNN extract the each and every portion of input image, which is known as receptive field. It assigns weights for each neuron based on the significant role of the receptive field. So that it can discriminate the importance of neurons from one another.

# LAYERS OF CNN

Convolutional neural networks are distinguished from other neural networks by their superior performance with image, speech, or audio signal inputs. They have three main types of layers, which are:

- Convolutional layer
- Pooling layer
- Fully-connected (FC) layer

# CNN LAYERS



# VGG19

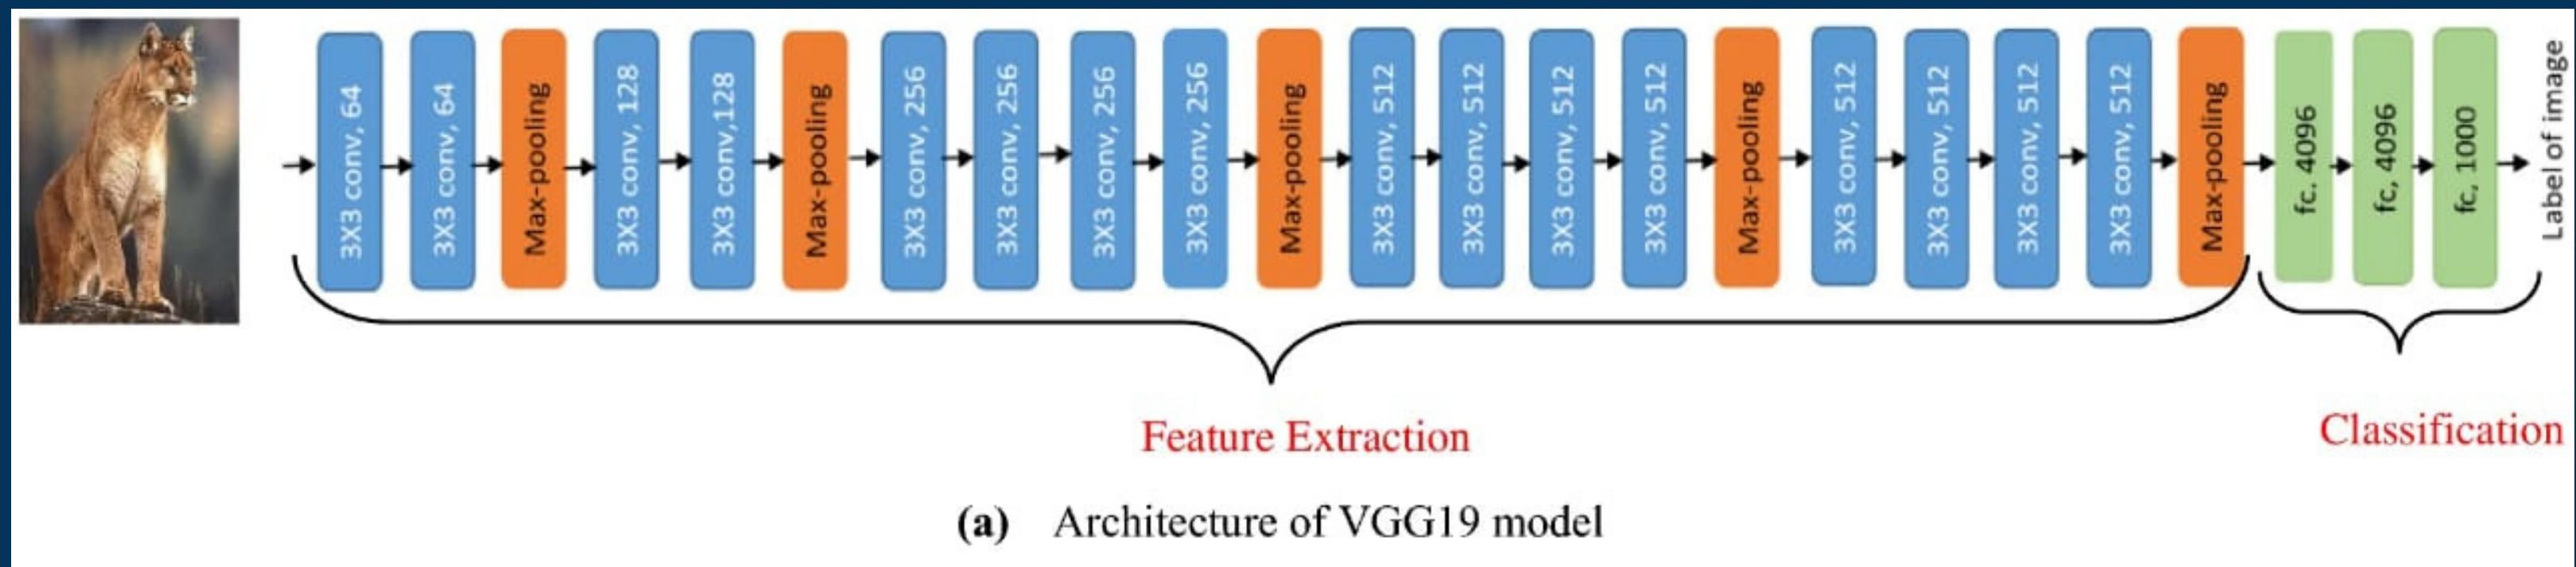
VGG19 proposed by Simonyan and Zisserman (2014) is a convolutional neural network that comprises 19 layers with 16 convolution layers and 3 fully connected to classify the images into 1000 object categories. We can load a pretrained version of the network trained on more than a million images from the ImageNet database. The pretrained network can classify images into 1000 object categories, such as keyboard, mouse, pencil, and many animals. As a result, the network has learned rich feature representations for a wide range of images.

- The network has an image input size of 224-by-224.

## VGG19

- So in simple language VGG is a deep CNN used to classify images. The layers in VGG19 model are as follows:
- Conv3x3 (64)
  - Conv3x3 (64)
  - MaxPool
  - Conv3x3 (128)
  - Conv3x3 (128)
  - MaxPool
  - Conv3x3 (256)
  - Conv3x3 (256)
  - Conv3x3 (256)
  - MaxPool
  - Conv3x3 (512)
  - Conv3x3 (512)
  - Conv3x3 (512)
  - MaxPool
  - Conv3x3 (512)
  - Conv3x3 (512)
  - Conv3x3 (512)
  - MaxPool
  - Fully Connected (4096)
  - Fully Connected (4096)
  - Fully Connected (1000)
  - SoftMax

# WORKING OF VGG19



# LIBRARIES USED

```
import os
import tensorflow as tf
os.environ['TFHUB_MODEL_LOAD_FORMAT'] = 'COMPRESSED'

import matplotlib.pyplot as plt
import matplotlib as mpl
mpl.rcParams['figure.figsize'] = (12, 12)
mpl.rcParams['axes.grid'] = False

import numpy as np
import PIL.Image
import time
import functools
from IPython.display import display
from IPython.display import clear_output
```

```
x = tf.keras.applications.vgg19.preprocess_input(content_image*255)
x = tf.image.resize(x, (224, 224))
vgg = tf.keras.applications.VGG19(include_top=True, weights='imagenet')
prediction_probabilities = vgg(x)
prediction_probabilities.shape
```

```
TensorShape([1, 1000])
```

```
predicted_top_5 = tf.keras.applications.vgg19.decode_predictions(prediction_probabilities.numpy())[0]
[(class_name, prob) for (number, class_name, prob) in predicted_top_5]

[('French_horn', 0.23891814),
 ('trench_coat', 0.106828496),
 ('bassoon', 0.06372758),
 ('sax', 0.052020073),
 ('suit', 0.03468096)]
```

```
vgg = tf.keras.applications.VGG19(include_top=False, weights='imagenet')

print()
for layer in vgg.layers:
    print(layer.name)

input_14
block1_conv1
block1_conv2
block1_pool
block2_conv1
block2_conv2
block2_pool
block3_conv1
block3_conv2
block3_conv3
block3_conv4
```

```
[101] content_layers = ['block5_conv2']

style_layers = ['block1_conv1',
                'block2_conv1',
                'block3_conv1',
                'block4_conv1',
                'block5_conv1']

num_content_layers = len(content_layers)
num_style_layers = len(style_layers)

[102] def vgg_layers(layer_names):
    # Load our model. Load pretrained VGG, trained on imagenet data
    vgg = tf.keras.applications.VGG19(include_top=False, weights='imagenet')
    vgg.trainable = False

    outputs = [vgg.get_layer(name).output for name in layer_names]

    model = tf.keras.Model([vgg.input], outputs)
    return model
```

```
style_extractor = vgg_layers(style_layers)
style_outputs = style_extractor(style_image*255)

#Look at the statistics of each layer's output
for name, output in zip(style_layers, style_outputs):
    print(name)
    print("  shape: ", output.numpy().shape)
    print("  min: ", output.numpy().min())
    print("  max: ", output.numpy().max())
    print("  mean: ", output.numpy().mean())
    print()
```

block1\_conv1

shape: (1, 512, 512, 64)  
min: 0.0  
max: 872.6996  
mean: 39.810673

block2\_conv1

shape: (1, 256, 256, 128)  
min: 0.0  
max: 4270.8975  
mean: 228.66785

```
def gram_matrix(input_tensor):
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensor, input_tensor)
    input_shape = tf.shape(input_tensor)
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)
    return result/(num_locations)
```

```
class StyleContentModel(tf.keras.models.Model):
    def __init__(self, style_layers, content_layers):
        super(StyleContentModel, self).__init__()
        self.vgg = vgg_layers(style_layers + content_layers)
        self.style_layers = style_layers
        self.content_layers = content_layers
        self.num_style_layers = len(style_layers)
        self.vgg.trainable = False

    def call(self, inputs):
        "Expects float input in [0,1]"
        inputs = inputs*255.0
        preprocessed_input = tf.keras.applications.vgg19.preprocess_input(inputs)
        outputs = self.vgg(preprocessed_input)
        style_outputs, content_outputs = (outputs[:self.num_style_layers],
                                         outputs[self.num_style_layers:])

        style_outputs = [gram_matrix(style_output)
                        for style_output in style_outputs]

        content_dict = {content_name: value
                        for content_name, value
                        in zip(self.content_layers, content_outputs)}

        style_dict = {style_name: value
                      for style_name, value
                      in zip(self.style_layers, style_outputs)}

        return {'content': content_dict, 'style': style_dict}
```

```
extractor = StyleContentModel(style_layers, content_layers)

results = extractor(tf.constant(content_image))

print('Styles:')
for name, output in sorted(results['style'].items()):
    print(" ", name)
    print("   shape: ", output.numpy().shape)
    print("   min: ", output.numpy().min())
    print("   max: ", output.numpy().max())
    print("   mean: ", output.numpy().mean())
    print()

print("Contents:")
for name, output in sorted(results['content'].items()):
    print(" ", name)
    print("   shape: ", output.numpy().shape)
    print("   min: ", output.numpy().min())
    print("   max: ", output.numpy().max())
    print("   mean: ", output.numpy().mean())

styles:
  block1_conv1
    shape: (1, 64, 64)
    min: 0.0001758667
    max: 37996.324
    mean: 506.44077
```

```
style_targets = extractor(style_image)['style']
content_targets = extractor(content_image)['content']

image = tf.Variable(content_image)

def clip_0_1(image):
    return tf.clip_by_value(image, clip_value_min=0.0, clip_value_max=1.0)

opt = tf.optimizers.Adam(learning_rate=0.02, beta_1=0.99, epsilon=1e-1)

style_weight=1e-3
content_weight=1e6
```

```
def style_content_loss(outputs):
    style_outputs = outputs['style']
    content_outputs = outputs['content']
    style_loss = tf.add_n([tf.reduce_mean((style_outputs[name]-style_targets[name])**2)
                          for name in style_outputs.keys()])
    style_loss *= style_weight / num_style_layers

    content_loss = tf.add_n([tf.reduce_mean((content_outputs[name]-content_targets[name])**2)
                           for name in content_outputs.keys()])
    content_loss *= content_weight / num_content_layers
    loss = style_loss + content_loss
    return loss
```

```
@tf.function()
def train_step(image):
    with tf.GradientTape() as tape:
        outputs = extractor(image)
        loss = style_content_loss(outputs)

    grad = tape.gradient(loss, image)
    opt.apply_gradients([(grad, image)])
    image.assign(clip_0_1(image))
```

```
train_step(image)
train_step(image)
train_step(image)
tensor_image(image)
```



```
import time
start = time.time()

epochs = 10
steps_per_epoch = 100

step = 0
for n in range(epochs):
    for m in range(steps_per_epoch):
        step += 1
        train_step(image)
        print(".", end='', flush=True)
        clear_output(wait=True)
        display(tensor_image(image))
        print("Train step: {}".format(step))

end = time.time()
print("Total time: {:.1f}".format(end-start))
```



```
def high_pass_x_y(image):
    x_var = image[:, :, 1:, :] - image[:, :, :-1, :]
    y_var = image[:, 1:, :, :] - image[:, :-1, :, :]

    return x_var, y_var

x_deltas, y_deltas = high_pass_x_y(content_image)

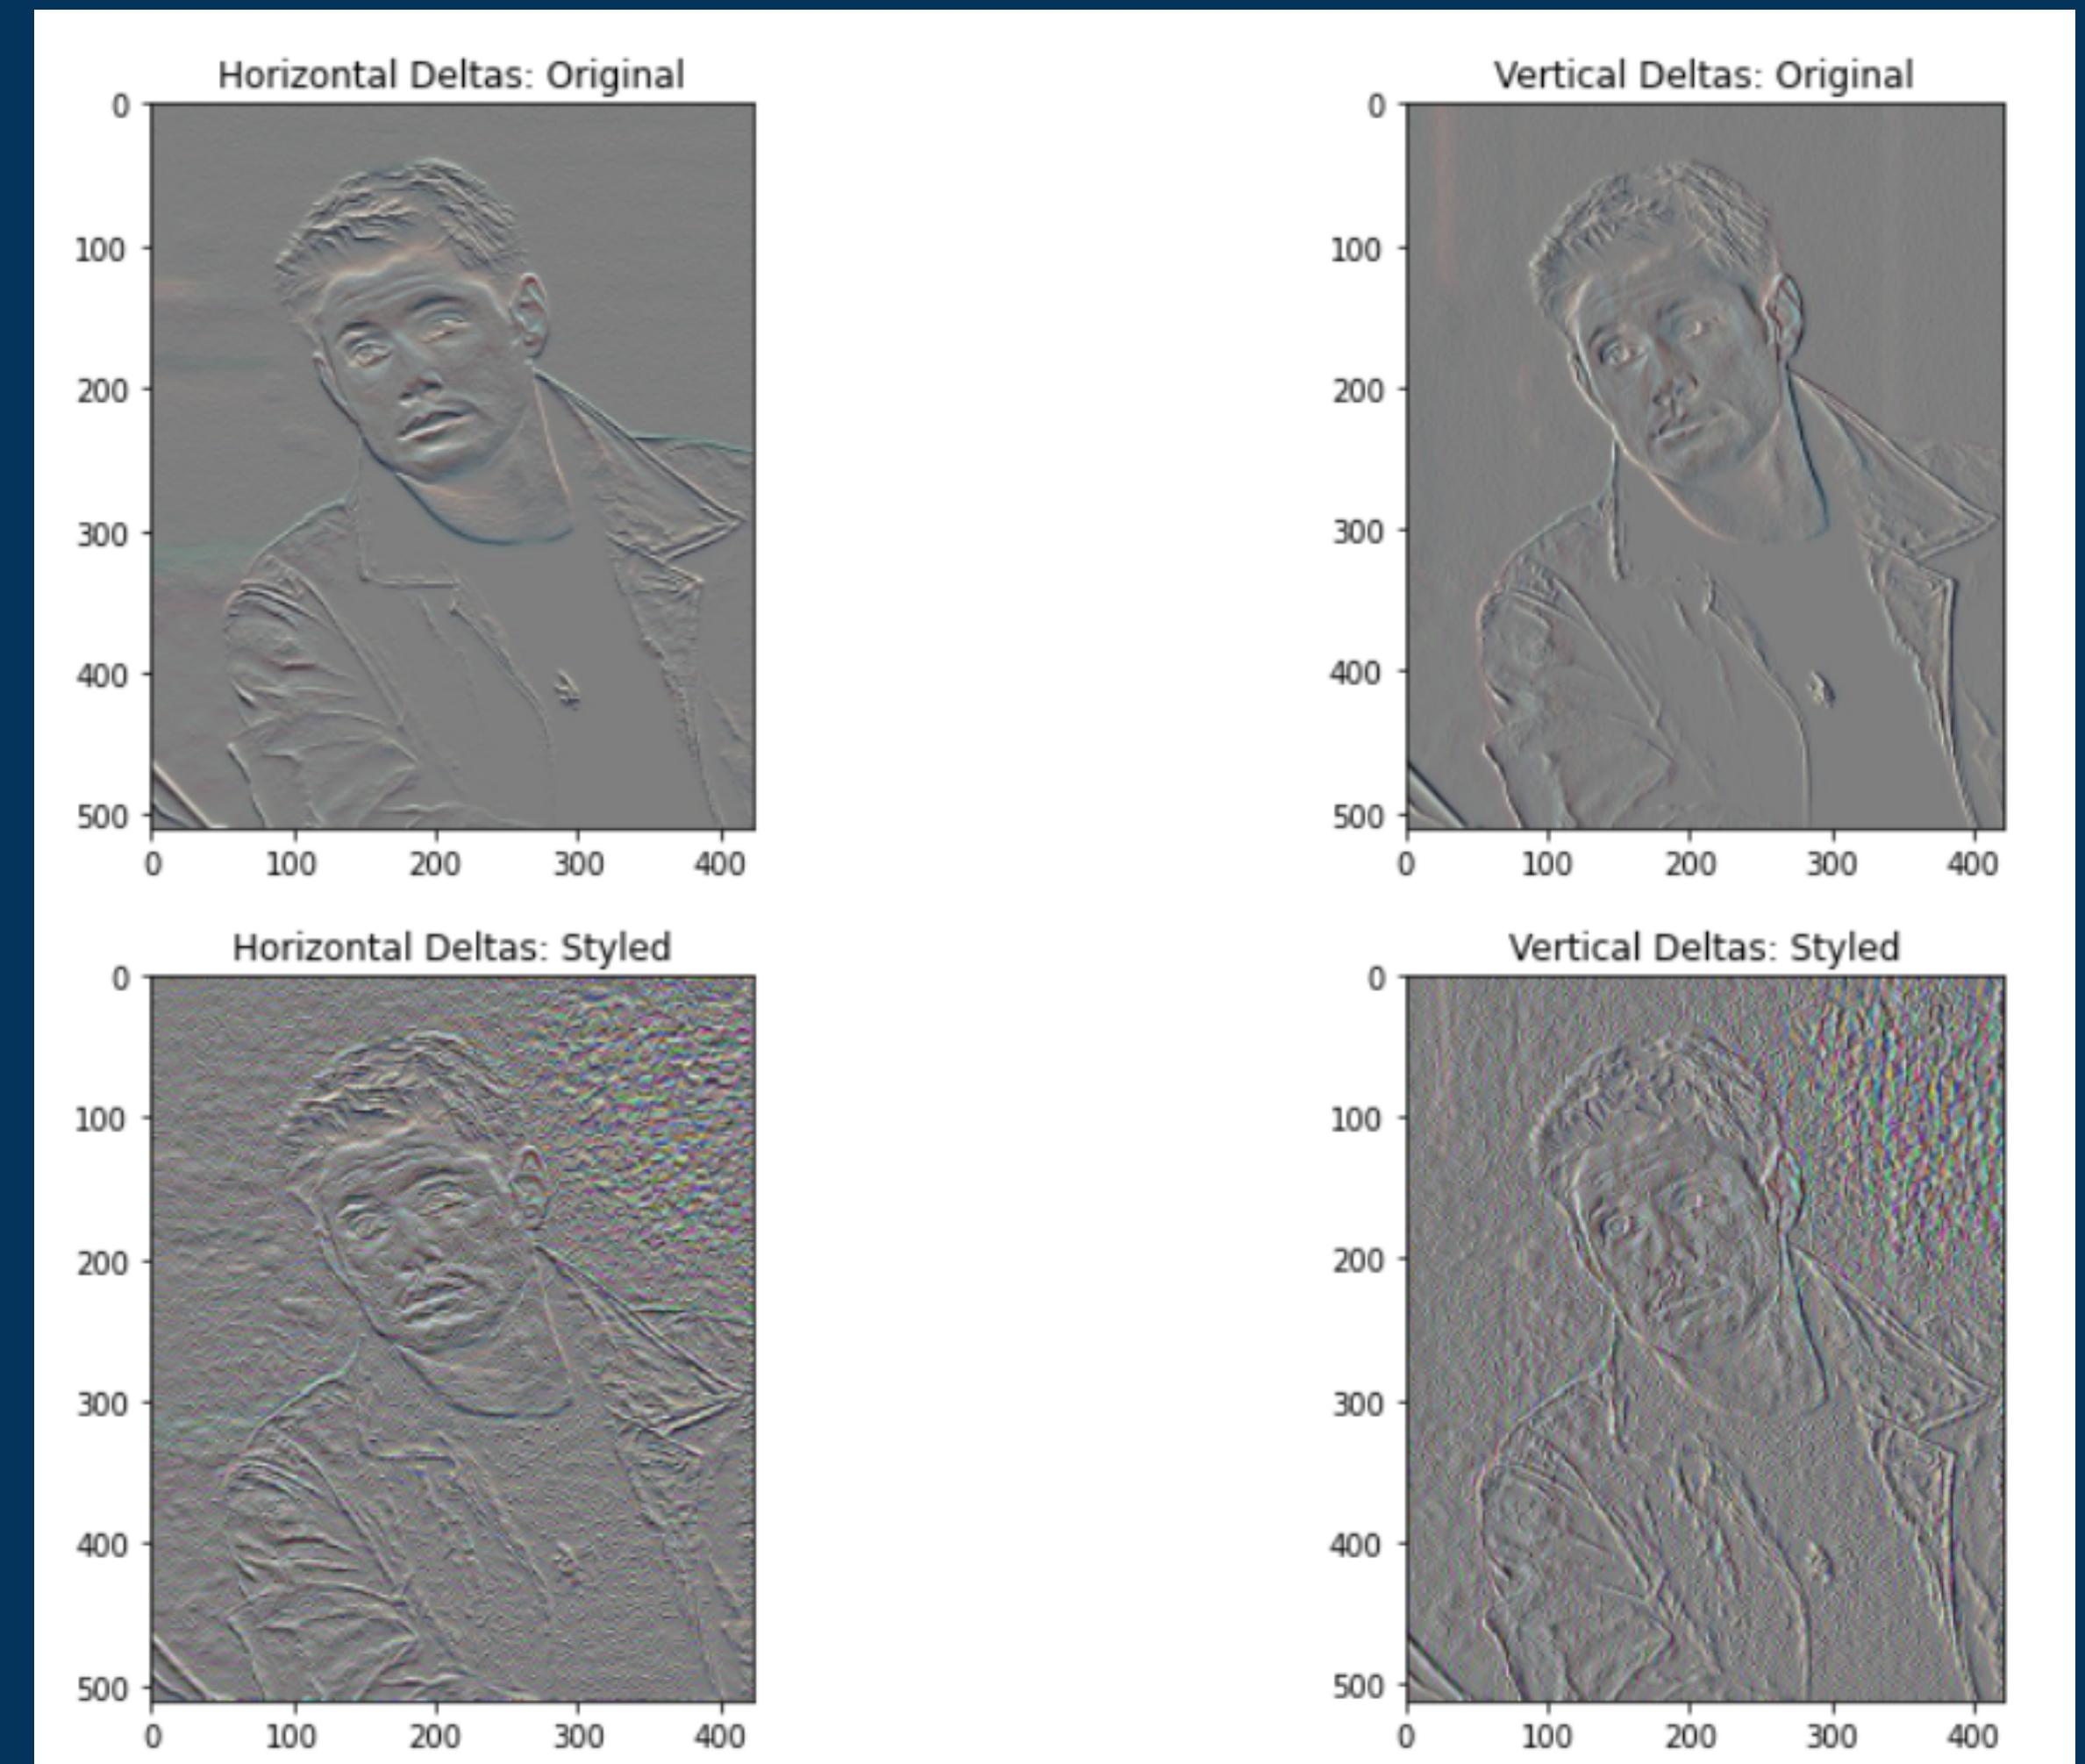
plt.figure(figsize=(14, 10))
plt.subplot(2, 2, 1)
imshow(clip_0_1(2*y_deltas+0.5), "Horizontal Deltas: Original")

plt.subplot(2, 2, 2)
imshow(clip_0_1(2*x_deltas+0.5), "Vertical Deltas: Original")

x_deltas, y_deltas = high_pass_x_y(image)

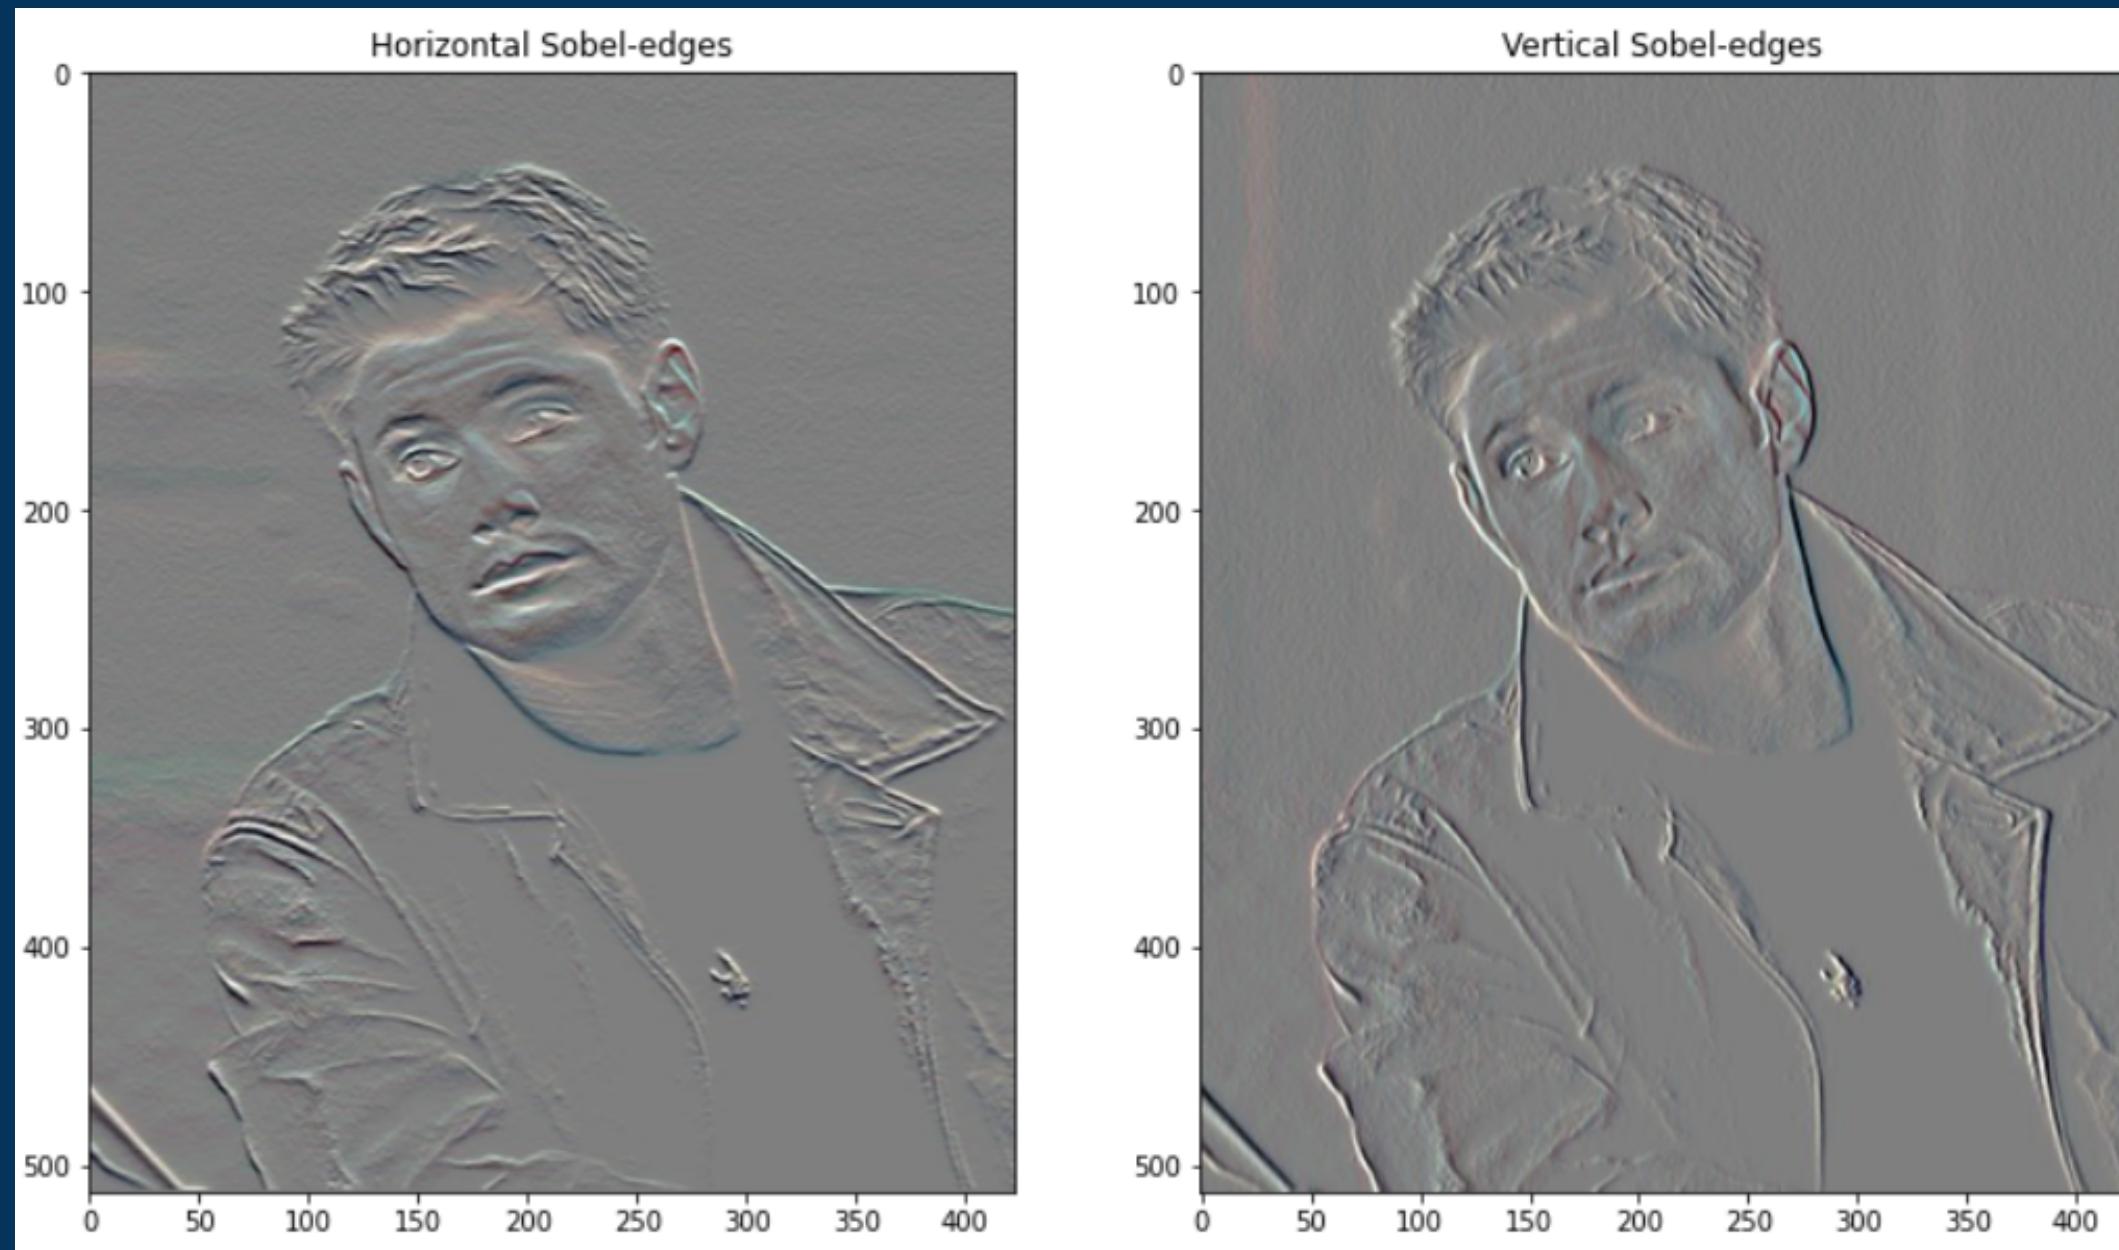
plt.subplot(2, 2, 3)
imshow(clip_0_1(2*y_deltas+0.5), "Horizontal Deltas: Styled")

plt.subplot(2, 2, 4)
imshow(clip_0_1(2*x_deltas+0.5), "Vertical Deltas: Styled")
```



```
plt.figure(figsize=(14, 10))

sobel = tf.image.sobel_edges(content_image)
plt.subplot(1, 2, 1)
imshow(clip_0_1(sobel[..., 0]/4+0.5), "Horizontal Sobel-edges")
plt.subplot(1, 2, 2)
imshow(clip_0_1(sobel[..., 1]/4+0.5), "Vertical Sobel-edges")
```



```
def total_variation_loss(image):
    x_deltas, y_deltas = high_pass_x_y(image)
    return tf.reduce_sum(tf.abs(x_deltas)) + tf.reduce_sum(tf.abs(y_deltas))
```

```
total_variation_loss(image).numpy()
```

```
142813.94
```

```
tf.image.total_variation(image).numpy()
```

```
array([142813.94], dtype=float32)
```

```
total_variation_weight=30
```

```
@tf.function()
def train_step(image):
    with tf.GradientTape() as tape:
        outputs = extractor(image)
        loss = style_content_loss(outputs)
        loss += total_variation_weight*tf.image.total_variation(image)

    grad = tape.gradient(loss, image)
    opt.apply_gradients([(grad, image)])
    image.assign(clip_0_1(image))

image = tf.Variable(content_image)

epochs = 10
steps_per_epoch = 100

step = 0
for n in range(epochs):
    for m in range(steps_per_epoch):
        step += 1
        train_step(image)
        print(".", end=' ', flush=True)
        clear_output(wait=True)
        display(tensor_image(image))
        print("Train step: {}".format(step))
```



```
file_name = 'stylized-image.png'
tensor_image(image).save(file_name)

try:
    from google.colab import files
except ImportError:
    pass
else:
    files.download(file_name)
```

# RESULT WITH VGG19 MODEL



Content Image

+



Style Image

=



Stylised Image

# RESULT WITH TF HUB MODEL

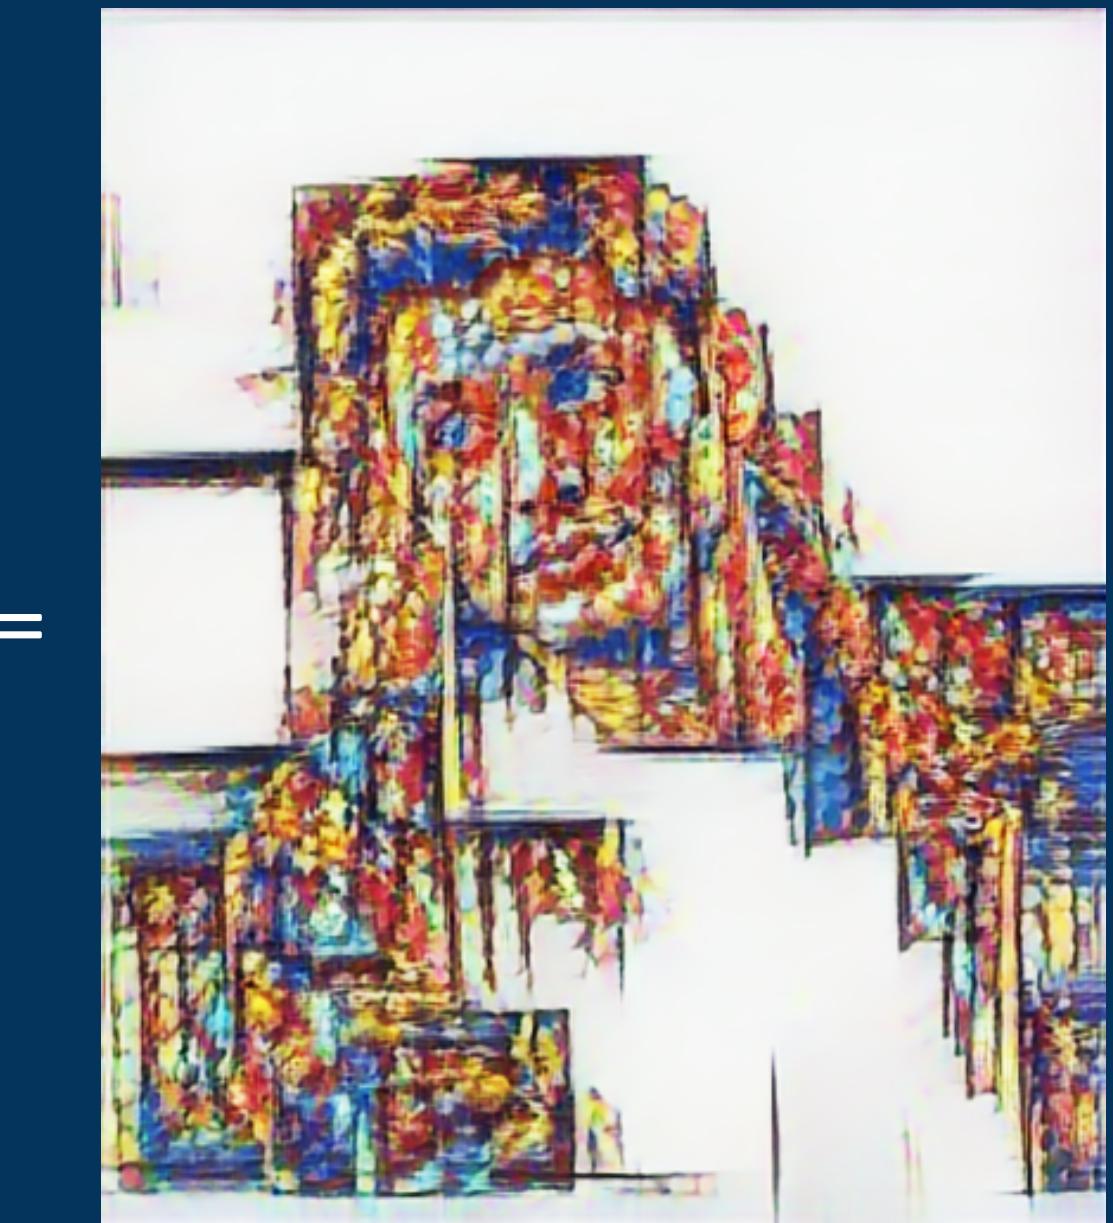


Content Image

+



Style Image



Stylised Image

# RESULT WITH VGG19 MODEL



Content Image

+



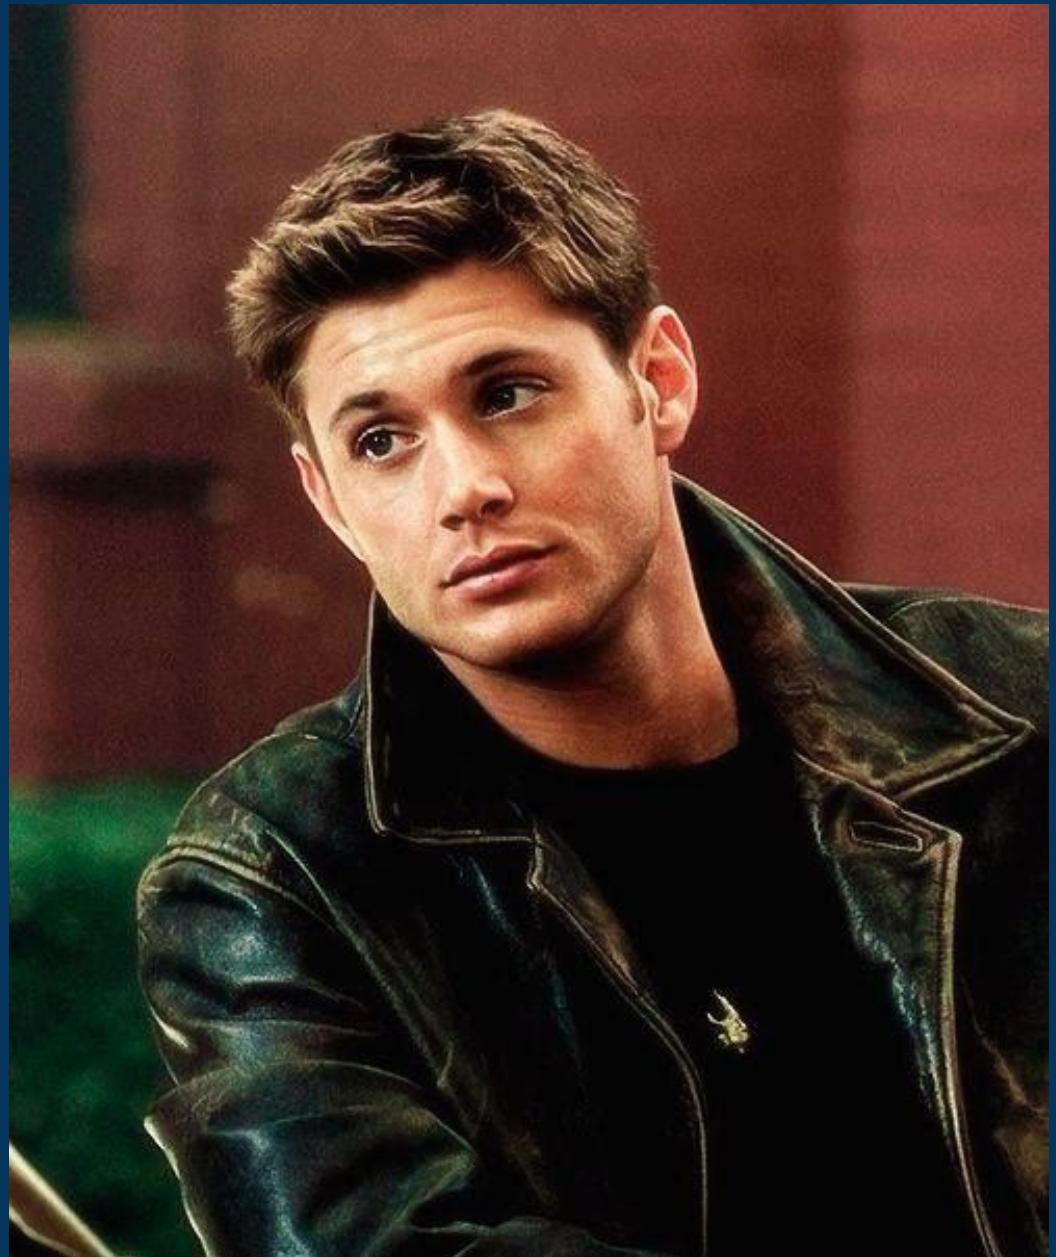
Style Image

=



Stylised Image

# RESULT WITH TF HUB MODEL



Content Image

+



Style Image

=



Stylised Image

The background features a minimalist design with large, semi-transparent teal circles of varying sizes scattered across a dark teal gradient background.

Design is the silent  
ambassador of your brand.